

Copyright
by
Venkata Subrahmanyam Govindarajan
2024

The Dissertation Committee for Venkata Subrahmanyam Govindarajan
certifies that this is the approved version of the following dissertation:

Modeling Intergroup Bias in Online Conversation

Committee:

Junyi Jessy Li, Supervisor

David Beaver, Co-supervisor

Kyle Mahowald

Malihe Alikhani

Modeling Intergroup Bias in Online Conversation

by

Venkata Subrahmanyam Govindarajan

Dissertation

Presented to the Faculty of the Graduate School of

The University of Texas at Austin

in Partial Fulfillment

of the Requirements

for the Degree of

Doctor of Philosophy

The University of Texas at Austin

May 2024

Dedication

To my mother and father.

Epigraph

Machine learning is like money laundering for bias. It's a clean, mathematical apparatus that gives the status quo the aura of logical inevitability. The numbers don't lie.

Maciej Ceglowski, [The Moral Economy of Tech](#)

Acknowledgments

This dissertation would not have been possible without the support of many people. I would like to thank my advisors Jessy, David, Kyle and Malihe, for their generosity over the last five years. They have always set aside their valuable time to talk to me about my progress, successes and failures. Thank you deeply.

I owe a debt of gratitude to Matianyu (Yuki) Zang, my collaborator over the last year, whose work on data curation and preprocessing has been invaluable.

My favorite activity in my PhD has been the research seminars, particularly the Computational Linguistics seminar, that I have attended weekly since Fall 2019. The conversations and arguments over research papers and projects with Katrin, Gabriella, Yejin, Juan Diego, Yating, Will, Hongli, Kanishka and many other researchers have taught me what it means to do research and be an academic, and made me a better one for it. Thank you for your feedback and thoughts on all my talks over the years.

The last five years have been easier to bear with a close group of friends — Benny, Gus, Hunter, and Ethan who I co-habited with through the beginning of the pandemic, and many years into it. Blake, Bryan and Erin for all those nights spent eating, drinking, at trivia, and commiserating as we navigated grad school. Sam, for choosing me as his side-kick with SXSemantics, and for being there for me when I needed community the most. Karoline, for a loving relationship that I will always cherish. Thank you all for welcoming me in, and for accepting all of me.

Everything in the end comes back to my mother and father, who have wished nothing else but for me to study and learn, and built opportunities for me that were never afforded to them. I am indebted to them for the opportunity to spend five years trying to understand this world a little bit more.

Abstract

Modeling Intergroup Bias in Online Conversation

Venkata Subrahmanyam Govindarajan, PhD
The University of Texas at Austin, 2024

SUPERVISORS: Junyi Jessy Li, David Beaver

Social bias in language is generally studied by identifying undesirable language use towards a specific demographic group, but we can enrich our understanding of communication by re-framing bias as *differences in behavior situated in social relationships* — specifically, the intergroup relationship between the speaker and target reference of an utterance. This dissertation aims to understand the nature of systematic variation between in-group and out-group speech — the intergroup bias. Chapter 2 describes systematic interactions between intergroup bias and interpersonal emotion, and finds that learned neural models can learn to classify the intergroup relationship in tweets with information about the speaker and target masked, outperforming expert annotators. To decipher what human-interpretable linguistic features are learned by these models, Chapter 3 describes probing experiments to understand if two pragmatic features — affect and specificity — are *used* by these models in differentiating in-group from out-group utterances. While affect and specificity have directional impacts on model prediction that align with our intuitions (positive affect utterances are more likely to be in-group), we found no interaction between these two variables, as we hypothesized generalizing from the Linguistic Intergroup Bias. Experiments up until this point also demonstrated the need for interpersonal language use *grounded* in non-linguistic world-state. To address this, Chapter 4 in-

investigates online comments from live game threads on forums (subreddits) dedicated to individual NFL teams. We find systematic linear relationships between the *form of referent* used to describe in-group and out-group teams, and the win probability for the in-group (game state) at the time of utterance. The better the state-of-the-world for the in-group, fans refer to the in-group less and the out-group more. Fans are also more likely to abstract away from events of the game and not refer to the in-group or out-group at all, choosing instead to express excitement, the more likely the in-group is to win. State-of-the-world as expressed in the win probability for the in-group turns out to be a well calibrated metric for intergroup bias, as evidenced by the linear relationships between several referent forms and win probability for the in-group. Data-driven analysis and experimentation with modern NLP tools in this dissertation thus revealed a form of intergroup bias ignored in previous literature, and outlines a promising picture of utilizing statistical information processing paired with careful data curation and analysis to understand subtle changes in human social behavior

Table of Contents

List of Tables	10
List of Figures	11
Chapter 1: Introducing intergroup bias	12
1.1 The Linguistic Intergroup Bias	13
1.2 Outline	15
Chapter 2: Political tweets and the intergroup bias	18
2.1 How to model intergroup bias?	20
2.2 Preliminary Analysis	25
2.3 Modeling the intergroup bias	29
2.4 Discussion & Conclusion	35
Chapter 3: Intergroup bias gets a counterfactual probe	37
3.1 Prior Work	38
3.2 Probing for the intergroup bias	40
3.3 Discussion & Conclusion	49
Chapter 4: Intergroup bias in Reddit NFL game comments	51
4.1 A new dataset of interpersonal language	54
4.2 Tagging & Annotation	58
4.3 Modeling the Intergroup bias with LLMs	69
4.4 Large-scale analysis of model tagged comments	76
4.5 Discussion & Conclusion	81
Chapter 5: Summary	84
5.1 Future Work	85
References	87

List of Tables

1.1	Predicted language variation in the LIB.	13
2.1	Tweets with in/out group and interpersonal emotion labels	22
2.2	Distribution of emotions in train-dev-test split	25
2.3	Proportion of emotions in different interpersonal contexts	26
2.4	Results on test set, with SD in parentheses, for intergroup relationship prediction task.	32
2.5	Top unigram and bigram features from NB-SVM model for each class.	33
2.6	F1 scores on test set, with SD in parentheses, for interpersonal emotion labelling task.	34
2.7	F1 scores on test set, SD in parentheses on out-group tweets. * indicates statistical significance ($p < 0.05$)	35
3.1	Predicted language variation in our more general formulation, using specificity and affect	39
3.2	Distribution of affect in train-dev-test split	42
4.1	Summary statistics of our dataset. Game comments are judged to be game thread comments that happened <i>during</i> the course of the corresponding NFL game. Comment timestamps were compared with publicly available start and end times of games.	57
4.2	Comments from the Chiefs and Eagles fans with the WP for the Chiefs in the middle. The WP for the Eagles is 1-WP for the Chiefs.	59
4.3	Summary statistics of expert annotated gold test set.	63
4.4	micro-F1 score (higher is better) over all tags, WER (lower is better) and GLEU scores (higher is better) on test split for all models.	75
4.5	Recall scores for each tag on test split for all models.	76
4.6	Table of slopes of feature of interest against increasing WP, alongside the r-squared showing how much of the variance is explained by the linear regression fit.	81

List of Figures

2.1	Emotions ordered by the number of examples where at least one rater uses a particular label. The color indicates the average interrater correlation.	24
2.2	Co-occurrence of emotions in our dataset.	27
2.3	Distribution of interpersonal emotions in unsupervised representations of tweets in our dataset. Orange indicates the emotion was present for that tweet. Each point represents one tweet from our dataset.	28
3.1	Top-100 LM Accuracy on train plus validation split of our dataset for different intervention strengths.	43
3.2	Flowchart describing the specificity intervention experiment and expected results.	45
3.3	% of test set classified as in-group plotted against number of INLP interventions for affect.	47
3.4	% of test set classified as in-group plotted against number of INLP interventions for specificity.	48
4.1	Frequency of any-group, in-group and null references in comments that fall in 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show some noisy trends. . .	67
4.2	Frequency of any-group and null references over all 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show clear linear trends.	77
4.3	Normalized frequency of references to the in-group, in-group using <i>we/us</i> and out-group over all 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show clear linear trends.	78
4.4	Normalized frequency of references to the in-group and out-group solely using names or <i>they/them</i> over all 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show clear linear trends.	79

Chapter 1: Introducing intergroup bias

Currently, most work studying bias in NLP situates bias as negative or pejorative language use towards an individual or group based on traits like race, gender, etc (Kane-ko and Bollegala, 2019; Sheng et al., 2019; Sap et al., 2020; Webson et al., 2020; Pryzant et al., 2020; Sheng et al., 2020). While these approaches greatly advance our understanding of bias in language and its impact and mitigation in NLP, focusing on specific demographic dimensions or an individual’s intent is limiting and not always practical. Research in psychology and social science suggests a different perspective. Bias can be seen as a relationship between people and groups, situated in context (Van Dijk, 2009); as such, bias refers to differences in behavior (in this case language use) as a result of differences in the relationship between speaker and target. The language we produce is biased in one way or another, whether we intend to or not, and whether that bias is positive, negative, or not clearly associated with any valuation (Beaver and Stanley, 2018).

So how do we study bias from this perspective? By accepting that the language we produce is biased, and is dependent on our social identity and relationships to the people we talk about (and to), we can focus on various aspects of modeling the subtle influence of social identity relationships on our linguistic production, as well as how language shapes and informs social identity in turn (Eckert, 2012). One such form of social meaning (Hall-Lew et al., 2021; Beltrama, 2020), that will be the focus of this dissertation is **intergroup bias**. The next section introduces the psycho-linguistic literature behind studies of the Linguistic Intergroup Bias, and its deficiencies that makes drawing inferences on social communication more broadly hard. By focusing on data-driven studies of in-group versus out-group speech in the wild, this dissertation aspires holistic understanding of how intergroup social differences has a bearing on the form, and meaning, of the language we produce.

	In-group	Out-group
socially desirable	abstract	concrete
socially undesirable	concrete	abstract

Table 1.1: Predicted language variation in the LIB.

1.1 THE LINGUISTIC INTERGROUP BIAS

The Linguistic Intergroup Bias hypothesis tries to explain why stereotypes persist, and how they are transmitted sub-consciously in daily conversation. In an intergroup context, and focusing on actions that are considered stereotypical for a group, socially desirable in-group behaviors and undesirable out-group behaviors are encoded at a higher level of **abstraction**, whereas socially undesirable in-group behaviors and desirable in-group behaviors are encoded at a lower level of abstraction (described in Table 1.1. A crucial underpinning of the theory is the Linguistic Category Model ([Semin and Fiedler, 1988](#)).

The Linguistic Category Model classifies predicates (words that can be used to describe people; like adjectives and verbs) on a scale of increasing abstraction – from verbs that are most concrete, to verbs that are less concrete, to adjectives that are most abstract. Abstract words (and thus statements) are taken to imply greater temporal stability and revealing of the character of a referent than their concrete counterparts. This provides an explanation for the persistence of bias and stereotypes through LIB – people tend to use abstract statements to talk about desirable in-group and undesirable out-group behaviors since they are **potentially more informative of the referent** (that is, indicative of future behavior). Consider the following utterances regarding a subject *Johnson* (examples taken from [Gorham \(2006\)](#)):

- (1)
 - a. The man police want to talk to probably hit the victims.
 - b. The man police want to talk to probably hurt the victims.
 - c. The man police want to talk to probably hated the victims.

- d. The man police want to talk to is probably violent .

hurt in (1-a) is a direct action verb; *hurt* in (1-b) is an interpretive action verb; *hated* in (1-c) is a stative and; *violent* is an adjective. Moving from (1-a) to (1-d), one can see how the information about the subject increases, while the information regarding a specific situation *decreases*. Thus, the *abstractness* of predicates increases from (1-a) to (1-d) according to the LCM.

The LIB uses the LCM ladder of abstraction to predict linguistic behavior: a speaker is more likely to describe an *out-group* individual with abstract predicates if the actions of the individual are socially undesirable, or negative stereotype congruent. Thus, white participants in the study from [Gorham \(2006\)](#) were more likely to describe the person whose picture they saw in a news report (whose race was varied as the experimental condition) using (1-c) or (1-d) if they were black (thus making them out-group), since it reinforces their negative stereotypes of African Americans. The converse holds for in-group referential utterances — white participants are more likely to use (1-a) or (1-b) to describe the person from the news report if he is presented as white.

There has been a wealth of work in psychology and psycholinguistics reproducing LIB in various domains such as crime reports and racial bias ([Gorham, 2006](#)); political news and party bias (?); as well as work exploring how LIB interacts with a speaker's prejudicial attitudes ([Schnake and Ruscher, 1998](#); [Greenwald and Krieger](#)). The LIB's strengths lie in the simplicity of its predictions, succinctly described in Table 1.1, and its focus on *abstraction* as a language feature — offering an attractive tie-in to cognitive mechanisms underlying prejudice and stereotypes. However, its weaknesses, closely tied to its strengths, also prevent it from being used to draw inferences of social communicative behavior at scale.

The LCM, upon which the LIB rests, while useful as an analysis of linguistic abstraction, suffers from a few drawbacks. The distinction between some of the classes, say

DAV and IAV, are not very linguistically motivated. DAVs are said to refer to ‘objective descriptions of observable behaviors’, all usages of that verb sharing a *physically invariant* component, while IAVs are said to refer to a general class of behaviors with positive or negative connotations. It remains to be investigated whether these definitions refer to something real (are DAVs less polysemous than IAVs?), but it would appear that a scale of abstractness would be more suited to this task. Some DAVs are more connotative than others (*hit* in (1-a) versus *perform*), whereas even within adjectives, some (like *athletic*) are more concrete than others (*loyal*).

Furthermore, the LCM constructs abstraction as simply a function of the verb/predicate. This is inherently limiting — can the subtle intergroup biases not be reflected in other parts of the utterance, or in the utterance as a whole? The simplicity of the LIB formulation is compounded by the ad-hoc nature of the defined axes of variation, especially the social desirability angle.

The LIB is a useful framework for analysis of utterances under very specific conditions — a focus on eliciting utterances from participants in constrained experiments, hand-coding social desirability as well as abstractness of predicates, and a focus on attitudes that are considered stereotypical of groups at the time. Real-world utterances about, or directed at, other people/groups show much more variation and diversity — does intergroup bias systematically influence real-world language use? This thesis is my program towards answering this question; To **characterize intergroup bias in real-world utterances through data curation, analysis and computational modeling**.

1.2 OUTLINE

This dissertation concerns intergroup bias in online conversation, which I aim to understand and study through large-scale data analysis and modeling. Chapter 2 introduces the notion of intergroup bias, and motivates why we need to study it in addition

to demographically-defined social biases. It defines various terms and concepts, and describes our first dataset of focus — tweets by US Congress members directed at other members. We find intrinsic statistical relationships between emotion and intergroup bias, which can further be learned and recognized by models for identifying if tweets are directed in-group or out-group with no knowledge of entities involved. This chapter was published as a paper at EACL 2023 (Govindarajan et al., 2023a).

Chapter 3 builds on top of the dataset and investigation in Chapter 1, and examines what *systematic, linguistic* changes can explain the differences between in-group and out-group speech. Extrapolating from the LIB, we define a new quadrant of intergroup language variation of the intergroup bias, operationalizing the ad-hoc axes of LIB towards automatically inferable, linguistically grounded variables. Through probing experiments, we discovered the limitations of a hypothesis driven approach towards *discovering* unknown, subtle intergroup variations in real-world language use, as well as the need for **grounding** our utterances in descriptions of events precipitating/preceding the utterance. While the hypothesis driven approach in this paper (published in the Findings of ACL 2024 (Govindarajan et al., 2023b)) failed, it was instructive and steered our focus to the issue of grounding and reference form itself.

Chapter 4 addresses this by introducing a new dataset of interpersonal utterances — over 6 million comments by NFL fans on live-game threads, grounding these interpersonal comments in events. Building upon the rich literature from the NFL statistics community, we utilize a real-valued number (the win probability) that succinctly describes the events preceding an utterance as it pertains most towards the intergroup bias — how well are things going for my in-group? We also introduce a novel way of modeling the intergroup bias, by tagging words in an utterance that refer to relevant entities as in-group or out-group.

Analyzing over 200,000 comments tagged with intergroup tags using the statistical information processing capabilities of modern Large Language Models (LLMs)

revealed a hidden variation not captured by the LIB: the **form of referent** when talking about the in-group or out-group changes systematically over time. Fans are more likely to abstract away from referencing a specific individual or team, towards a description of events in general, the more likely their team (the in-group) is to winning. Furthermore, this trend is remarkably linear over win probabilities. References to the out-group remain steady across win probabilities.

Overall, the findings in this thesis constitute the first data-driven large scale study of intergroup bias in real-world language use. The findings add much needed color and linguistic rigor to the LIB hypothesis. Future work needs to expand this work to more domains, to gain a holistic understanding of how social structures and relationships mediate our language use subconsciously.

Chapter 2: Political tweets and the intergroup bias

In psychological work on Linguistic Intergroup Bias (Maass, 1999), bias originates from the relationship between the speaker and target of an utterance, i.e. their **interpersonal dynamics**, and manifests later in subtle ways. Consider the utterances (tweets) in (1), drawn from our collected data in which the identity of the speaker and target are masked:

- (1) a. **In-group:** We stand w @Doe, who has seen a lot worse than cheap insults from an insecure bully. #MLKDAY weekend.
- b. **Out-group:** Parents and families live in constant fear for their children with food allergies. A worthy bipartisan cause - thank you @Doe for your leadership on this issue.

Both express support and admiration towards the target referent *Doe* – however, the second example uses words indicative that the speaker and target do not share a relevant social identity (in this case, their political party), expressed by words like *bipartisan*. The intensity of admiration expressed is also greater in (1-a) than (1-b). Thus, these two seemingly similar statements differ along interpersonal dimensions that are instructive as to how the bias of the speaker seeps into the utterance.

We now introduce two new tasks that directly model language use in terms of two interpersonal dimensions: (i) **intergroup relationship (IGR) prediction**, where we seek to understand how people talk about others who they consider to be in their same social group (in-group), versus those they consider outside their social group (out-group), and (ii) perceived **interpersonal emotion detection**, where we situate these differences in terms of the emotion expressed in text *towards or in connection with* a target individual described in the utterance. Note that *interpersonal* emotion is dif-

ferent from a more standard, utterance level emotion detection task, as illustrated in row 2 of Table 2.1 which has seemingly opposing emotions.

We present a first-of-its-kind, *annotated* dataset for fine-grained interpersonal emotion detection, consisting of 3,033 tweets from members of the US Congress; all of these tweets mention another Congress member, hence providing us with ‘found supervision’ for IGR prediction (whether the speaker and the target belong to the same political party). Our analyses show that while positive interpersonal emotions appear in both in- and out-group situations, negative emotions like anger and disgust are overwhelmingly present in the latter. Meanwhile, human judgments for in vs. out-group membership on this dataset are overly reliant on the polarity of emotion; specifically, human judges are much *less* likely to attribute positive emotions towards out-group targets.

Baseline performances for perceived interpersonal emotion detection shows that this is a challenging task, as is consistent with existing work in emotion detection in general [Demszky et al. \(2020\)](#). In particular, emotions in this dataset are often expressed with considerable subtlety, likely a characteristic of official political speech. To investigate whether IGR and emotions are intertwined and useful towards each other, we further developed a multi-task model for the prediction of both. We found compelling evidence that multi-tasking IGR and interpersonal emotion improves performance on both tasks with over 10% improvement in detection of disgust in out-group contexts, and 3% improvement in IGR prediction.

To summarize the contributions of this chapter, we tackle **intergroup bias**, a notion of bias rooted in social psychology that applies to all the various differences in the ways that people talk about others in their in-group or out-group. Standard bias tasks in NLP, and the broader goal of debiasing models could thus be set in a more general context. We present the first dataset to study both intergroup membership and emotion, which allows us to analyze both human and model behavior in terms of how the two interact with each other.

2.1 HOW TO MODEL INTERGROUP BIAS?

2.1.1 TWO DIMENSIONS OF INTERGROUP BIAS

Our aim is to build a generalized, data-driven approach towards studying bias situated in **interpersonal utterances**, which we define as any utterance where there is a target individual being talked about or referred to. Our goal is to model two novel tasks described below; examples are shown in Table 2.1.

Intergroup relationship IGR is defined by the relationship between the speaker and target of an utterance. People belong to multiple social groups as part of their identity, however usually only some identities are salient in an utterance in context. We define *in-group* utterances as ones where the speaker and target are in the same social group, and *out-group* utterances as one where they are in different social groups. Given an utterance u written by an individual s with target t , the IGR prediction task classifies whether s and t belong to the same social group within the context of u .

Interpersonal Emotion We define *perceived* interpersonal emotion as the emotion expressed by a speaker s *towards, or in connection with* the target t of the utterance u , as perceived by a reader. We use the Plutchik wheel of emotions, which is widely adopted in the community, as the basis of our emotion taxonomy [Plutchik \(2001\)](#); we use the 8 fundamental emotions (*admiration, anger, disgust, fear, interest, joy, sadness, surprise*) instead of the full 24 emotions in the wheel due to data sparsity. Interpersonal emotions may be different, or a subset of, emotion for the whole of an utterance, as illustrated in rows 2, 3 and 4 of Table 2.1. Given an utterance u written by an individual s with target t , the interpersonal emotion detection task identifies the perceived emotion of s towards the target t .

2.1.2 DATA SOURCES AND PREPROCESSING

In our area of focus, we require natural language data which satisfies the following criteria: (1) Each utterance must have at least one target about whom the utterance mainly concerns. (2) The relationship between the speaker and the target must be inferred based on metadata or other information. Specifically, we are interested in aspects of their social identity that they share or differ on.

The dataset we collect comes from tweets by members of US Congress where other members are mentioned in the same tweet. We use this as a convenient testbed: each member’s group affiliation (i.e., their party identity) is public, thus we can easily know whether the speaker is tweeting to a target in their own party or not.¹ In other words, this dataset gives us “found supervision” for our first task of IGR prediction. For our second task, we annotate a subset of these tweets for perceived interpersonal emotion; this is, to our knowledge, the first dataset dedicated to interpersonal emotion.

Social media text like tweets offer a fertile ground for our study. A focus on tweets with *mentions* in them satisfies our first criterion – people generally use mentions to say something about or towards another individual on twitter. Tweets by members of US Congress are a matter of public record, and we can infer the social relationship (in terms of party affiliation) between speaker and target using publicly available information. We prioritize working with a dataset of tweets by members of US Congress (downloaded using the Twitter API) between 2010 and 2021, spanning two presidencies, during which both parties held power in Congress. We filter these tweets to exclude retweets, and include those tweets that mention *at most* one other member of Congress whose party affiliation is known. We believe these 2 assumptions are sufficient to arrive at a dataset of tweets where the speaker is talking towards/about *one* target. Thus, we restrict ourselves to two social groups in this sphere — Democrat

¹For simplicity, we do not consider other factors such as the home state of a congress member.

Tweet	Interpersonal Emotion	In/Out group?
As @Doe says, the times have found each and every one of us to Defend our Democracy For The People. Worth reading every line.	Admiration	In-group
Freedom has no greater nor tougher champion than @Doe. My prayers are with him and his family.	Admiration & Sadness	In-group
You don't get to decide what's "fine," @Doe. The constitution does. #DefendOurDemocracy #WednesdayThoughts	Anger & Disgust	Out-group
Thank you again Senator @Doe for leading the SRF WIN Act[...] I'm proud to be a co-sponsor	Admiration & Joy	Out-group

Table 2.1: Tweets with in/out group and interpersonal emotion labels

and Republican parties in the US. We sample an equal number of in-group and out-group tweets from a large sample consisting of all tweets by members of Congress. Apart from years 2010–2012 and 2021 which contained fewer tweets due to sparsity issues, we sampled at least 300 tweets each year.

2.1.3 INTERPERSONAL EMOTION ANNOTATION

While we can infer whether a tweet is in-group or out-group based on the identity of speaker and target whose political affiliations are known, we still require annotated data on perceived interpersonal emotions. Interpersonal emotions vary in subtle ways from sentiment or overall sentiment of utterances: an utterance can have negative sentiment overall, but still convey positive emotions towards the target of the sentence (expressing admiration at someone’s death for instance). For this reason, we devise an annotation schema for annotating *the emotion expressed by speaker s towards target t* .

Instructions Annotators are presented with a tweet, with the identity of the speaker unknown and that of the target masked with a placeholder name **@Doe** to minimize potential biases of the annotators’ prior knowledge of party affiliation intruding into the annotation:

(2) If **@Doe** can get her hair done in person, Congress can vote in person...

Annotators are instructed to read the tweet and select only the most notable emotion(s) they think are expressed by the tweet author *in connection with @Doe*. To aid annotators, we provide examples of the 8 Plutchik emotions (*joy, admiration, fear, surprise, sadness, disgust, anger and interest*) expressed as interpersonal emotions in tweets. Annotators are also shown a schematic of the Plutchik wheel of emotions, which acquaints them with how the emotions are related to one another in our framework. Annotators are allowed to select more than one emotion to account for emotion co-occurrence. We also explicitly tell annotators that more than one of the emotions can be present in the tweets, to encourage them to select all interpersonal emotions expressed. They are also allowed to not choose any emotion.

Annotation To obtain reliable annotations, we prequalify annotators using a qualifying task. Annotators were recruited on Mechanical Turk using a qualifying task where they were asked to annotate 6 tweets using the schema shown above. We restricted the qualification task to annotators living in the USA who had attempted at least 500 HITS and had a HIT approval rate $\geq 98\%$. After manual inspection, 6 annotators were qualified for bulk annotation. Each tweet was annotated by three different annotators. To ensure annotators were paid a fair wage of at least 10\$ an hour, we paid annotators \$0.50 per HIT. Each HIT involved annotating 3 tweets, which we estimate to take on average 3 minutes to complete. In total, 3,033 tweets between 2010 and 2021 were annotated with perceived interpersonal emotion.




Figure 2.1: Emotions ordered by the number of examples where at least one rater uses a particular label. The color indicates the average interrater correlation.

Agreement To measure agreement between annotators on the Plutchik-8 emotion wheel, we use the Plutchik Emotion Agreement (PEA) score from [Desai et al. \(2020\)](#). The PEA score addresses the issue of penalizing all disagreements equally, by penalizing dissimilar emotion annotations higher than more similar ones (according to the Plutchik wheel). Our PEA score is 0.73. The original PEA formulation used the best(max) pair of emotion annotations between two workers. Taking the *worst* combination of emotions between two workers (averaged over all tweets and workers), the PEA (min) score is 0.60. Overall, we find moderate to high agreement on fine-grained interpersonal emotions. In Figure 2.1 we also present interrater correlation, a metric used in [Demszky et al. \(2020\)](#); we see that distributions are similar.

Aggregation We consider a tweet to have a certain emotion label if at least 2 out of 3 annotators agree that the particular emotion was present in the tweet. A total of 638

Emotion	Train	Dev	Test
Admiration	467	64	58
Anger	225	40	46
Disgust	206	32	43
Fear	1	0	0
Interest	701	83	84
Joy	801	107	106
Sadness	72	11	11
Surprise	1	0	0
<i>No Emotion</i>	519	56	63

Table 2.2: Distribution of emotions in train-dev-test split

tweets have no interpersonal emotion associated with them. We employ a 80-10-10 train-dev-test split on our data.

The number of annotated examples (tweets) per emotion is shown in Table 3.2. We omit *fear* and *surprise* from future tables due to the absence of annotated examples.

2.2 PRELIMINARY ANALYSIS

How are emotions distributed? When observing the distribution of aggregated emotion labels themselves, a clear pattern emerges as seen in Table 2.3. Negative emotions such as anger and disgust are almost always expressed in out-group settings, while positive emotions are present in both in-group and out-group settings. A similar distribution of emotions was observed for Democrats and Republicans — members of both parties reserved their public anger and disgust for members of the other party. This reflects an innate bias in terms of the distribution of interpersonal emotions per situation, and warrants future work to explore negative interpersonal emotions in an in-group setting.

Figure 2.2 shows the co-occurrence of interpersonal emotions in our dataset. We can see that emotions that are farther apart and more dissimilar, such as admiration and

Emotion	All	In-Group	Out-Group
Admiration	15.5	22.2	9.1
Anger	8.2	1.0	15.1
Disgust	7.4	0.3	14.2
Interest	22.9	27.2	18.6
Joy	26.7	32.2	21.4
Sadness	2.5	2.6	2.4
<i>No Emotion</i>	16.8	14.5	19.1

Table 2.3: Proportion of emotions in different interpersonal contexts

disgust, joy and sadness, co-occur infrequently. Emotions that are closer such as anger and disgust, admiration and joy, co-occur much more often. The only outlier is the higher than normal co-occurrence of admiration with sadness — after a closer examination, this can be attributed to tweets expressing admiration and sadness at the passing, or end of the career, of a fellow congressperson.

Who were the targets of negative emotions? On further analysis, it appears that most of the out-group disgust and anger is directed at 3 handles – @speakerriyan, @speakerpelosi, and @speakerboehner who were all Speakers of the House of Representatives over most of the time period of our dataset. 63.7% of disgust and 64.3% of anger is directed towards these three twitter handles. 11.9% of all tweets in our dataset are directed at these handles, indicating the preponderance of negative interpersonal emotion directed at the Speaker of the house. However, we note that negative emotions like anger and disgust were still expressed towards 51 and 45 different individuals in our dataset, respectively.

Can humans predict in/out-group? While our data naturally comes with “gold” IGR labels, what is unexplored is whether the distinction between in-group and out-group speech is prominent and noticeable by humans. Additionally, it is also unclear if hu-




Figure 2.2: Co-occurrence of emotions in our dataset.

mans might have their own expectation of how in/out-group speech should be characterized.

Concretely, we investigate if human annotators were capable of accurately performing the IGR prediction task when the speaker and target are masked. Two collaborators, one a social science graduate student, and the other a computational linguistics graduate student, annotated 50 random tweets from our validation data which they had not been exposed to earlier for in/out group labels. Their Fleiss κ agreement score was 0.64, indicating moderate agreement.

To check how accurate their judgements were, we calculate for each annotator their F1 score against our “gold” in/out group labels. Their F1 scores on these 50 tweets were 0.67 and 0.63, which as we will discuss in Section 2.3.2, only match simple baselines of supervised systems. Annotators comments indicate that they overly relied on the sentiment of tweets to make the classification — positive sentiment means




Figure 2.3: Distribution of interpersonal emotions in unsupervised representations of tweets in our dataset. Orange indicates the emotion was present for that tweet. Each point represents one tweet from our dataset.

in-group and negative sentiment means out-group. While negative emotions are over-represented in out-group situations as Table 2.3 shows, our dataset contains a substantial presence of out-group tweets with positive interpersonal emotions as well. Annotators also noticed some lexical cues like ‘bipartisan’ that are indicative of out-group tweets.

Pre-trained representations capture interpersonal emotions Pre-trained language models have been found to learn sentence representations that cluster by domain without supervision (Aharoni and Goldberg, 2020). We wished to investigate if any of our annotated properties cluster inherently in reduced representations of the tweets in our data. We use unsupervised representations from BERTweet (Nguyen et al., 2020), a language model pre-trained on 850M English tweets. We take the 768 dimensional embeddings from the final layer of the <s> token in BERTweet, and dimen-

sionally reduce them to 2 dimensions using UMAP (Sainburg et al., 2021). Figure 2.3 shows the distribution of tweets, color coded for interpersonal emotions. While there is a lot of overlap between representations when stratified by emotion, we can see that some emotions that are intuitively opposite, like admiration & disgust, joy & sadness are moderately separable. This indicates that interpersonal emotions do define some topic or domain level properties of a tweet.

2.3 MODELING THE INTERGROUP BIAS

2.3.1 EXPERIMENTS

We detail our experiments for the two novel tasks discussed in Section 2.1.1: predicting the IGR (in-group or out-group) given a tweet, and predicting the interpersonal emotion given a tweet. We present baselines for the two tasks separately, and also present a multi-task model to gauge the extent to which knowledge of IGR may help in predicting interpersonal emotion, and vice versa.

2.3.1.1 Intergroup Relationship

Sentiment-Rule Our first baseline is a rule-based one leveraging coarse sentiment: if a tweet’s sentiment is predicted to be negative, classify it as out-group; if positive, classify it as in-group; and if neutral, classify it as either in-group or out-group randomly. We use a RoBERTa-Base model fine-tuned for sentiment on tweets (Barbieri et al., 2020) to extract the sentiment of each tweet in our dataset.

NB-SVM As a second baseline, we build an SVM model that uses Naive-Bayes log-counts ratios of unigrams and bigrams (Wang and Manning, 2012).

BERTweet We use BERTweet (Nguyen et al., 2020), a language model pre-trained on 850M English tweets as our dataset consists purely of English language tweets. A classification head is placed on top of the language model. We also experiment with a version where the language model parameters are frozen, and only the classification head parameters are fine-tuned (BERTweet-ft).

The input to all models is only the tweet with no other context, and the target masked with a placeholder @USER.

2.3.1.2 Interpersonal Emotion

EmoLex As a baseline model for interpersonal emotion identification, we rely on EmoLex (Mohammad and Turney, 2013). EmoLex consists of 14,182 crowdsourced words associated with the 8 basic Plutchik emotions. Critically, these words appear in emotional contexts, but are not necessarily emotion words themselves. EmoLex counts occurrences of words from its lexicon in an utterance, and assigns a normalized score for each emotion based on occurrence frequency. We consider an emotion to be on, if it’s normalized score is ≥ 0.001 . While EmoLex has issues with regards to its context insensitivity and the social biases built into its lexicon (Zad et al., 2021), we include it as a baseline to understand to what extent interpersonal emotions can be deduced using a lexicon.

BERTweet We use the same BERTweet model as earlier. We add a dense output layer on top of the pretrained model for the purposes of finetuning, with a sigmoid cross entropy loss function to support multi-label classification. The loss is weighted for each of the 8 emotion labels with the ratio of positive and negative examples to increase precision. If none of the 8 emotion labels are flipped on, we consider that to be the ‘No Emotion’ label, i.e. there is no interpersonal emotion between speaker and target in the tweet. We experiment with a version of the model where the language

model parameters are frozen and only the labelling head parameters are fine-tuned (BERTweet-ft).

2.3.1.3 Multi-Task Model

In § 2.2, we observed that the emotions anger and disgust are overwhelmingly present in out-group situations. Thus, we hypothesize that IGR information would be useful towards interpersonal emotion identification, and vice versa. To test this hypothesis we train a multi-task model. The model is trained to predict *both* the IGR label and emotion using shared parameter finetuning.

We use the same BERTweet model as earlier. We add two dense output layers on top of the pretrained model, one for classifying IGR and another for labelling interpersonal emotion. Both heads share the same parameters below. These are trained with same loss as earlier individual models. The model alternates between finetuning for group relationship and emotion over every training item.

2.3.1.4 Implementation

We use bertweet-base pretrained embeddings from the Huggingface model hub (Wolf et al., 2020). All models are fine-tuned for a maximum of 20 epochs with early stopping. Early-stopping patience for models trained on each task separately is 3. The patience for the multi-task model is set at 5 as the multi-tasking setup led to slower convergence. The learning rate for the classification heads was set at 5e-3 while the learning rate for the internal language model parameters was set at 2e-5. Dropout probabilities in classification heads was set at 0.1. The best performing model before early stopping on validation data was chosen in all cases. We report F1 scores averaged over 3 random restarts for all models, with the standard deviation in parentheses next to the mean.

Model	F1	Model	F1
Majority class	51.1	BERTweet	74.1 (3.3)
Sentiment-Rule	56.3	BERTweet-ft	66.5 (1.6)
NB-SVM	62.5	Multitask	77.3 (0.8)
Human	66.7		

Table 2.4: Results on test set, with SD in parentheses, for intergroup relationship prediction task.

2.3.2 RESULTS AND ANALYSES

Intergroup relationship In modeling IGR, we find that `Sentiment-Rule` performs not much better than chance (Table 2.4). This underscores one strength of our data, which contains a sizable number of out-group tweets with positive interpersonal emotion attached to them. The `NB-SVM` model based on unigrams and bigrams performs slightly better, and picks up on some obvious out-group lexical cues like the lemma ‘bipartisan’, as shown in Table 2.5. The `BERTweet` model performs substantially better, performing over 10 points better than humans. The model, with only the classification head fine-tuned, leaving the language model parameters intact(`BERTweet-ft`) performs about 10 points above chance — indicating that there may be features advantageous towards this task in the vanilla LM itself.

Interpersonal Emotion We find that the `EmoLex` baseline, which relies purely on lexical cues, performs dismally on our data, with poor performance in both in-group and out-group settings (Table 2.6). This is a strong indication that emotions are expressed more implicitly in this dataset. The `BERTweet` model performs substantially better, indicating that interpersonal emotions, even if implicit, can be learned.

Multitask Model As Table 2.4 shows, Multi-tasking the two tasks leads to a noticeable improvement in F1 for IGR prediction, with the differences being statistically

In-group	Out-group
thanks, love, count me birthday, my colleague	thanks, bipartisan, restore kind, resignation

Table 2.5: Top unigram and bigram features from NB-SVM model for each class.

significant using a bootstrap test ($p < 0.05$; [Berg-Kirkpatrick et al., 2012](#)); the multitask model is also more stable with much lower variance across runs. These results suggest that interpersonal emotion is useful towards IGR prediction.

Table 2.6 shows that the performance of the multitask model on predicting interpersonal emotions is significantly better than the BERTweet model ($p < 0.05$) on emotions like *disgust*, which suggests that IGR is useful towards the task of emotion identification. Furthermore, multitasking boosted performance at predicting the *no emotion* label by 20%. Table 2.7 compares the multitask model’s performance against the BERTweet model in *out-group* settings (where most of the gains were found) for 3 emotions — illustrating the boost in performance afforded by joint modeling of IGR and emotion for *disgust*. The 3 emotions listed also showed significant differences in their distribution in in-group and out-group settings.

Humans vs. Models Comparatively, we find that model performance exceeds human performance on the task of in-group versus out-group prediction, albeit not on the same dataset. The model’s main driver of performance is its high accuracy on positive intergroup emotion out-group tweets, such as those expressing admiration or joy. Human annotators consistently fall back on the heuristic that sentences with positive affect probably imply that the speaker is talking about someone in their in-group. But it is not the case in the political domain, where overtures to bipartisanship serve as useful signals. For instance, both (3-a) and (3-b) express admiration towards the target Doe, where the first is in-group while the second is out-group. The call

	EmoLex	BERTweet	BERTweet-ft	Multi-task
Admir.	37.5	70.3 (3.7)	40.7 (1.1)	68.9 (1.6)
Anger	26.6	71.3 (11.2)	23.0 (3.4)	69.3 (3.3)
Disgust	25.5	47.1 (21.6)	13.0 (4.1)	74.5 (7.1)
Interest	0	53.1 (3.3)	5.8 (2.4)	51.5 (8.5)
Joy	48.4	85.9 (1.9)	71.3 (1.4)	83.6 (1.3)
Sadness	4.3	11.1 (9.6)	0	33.6 (18.5)
<i>No Emotion</i>	22.2	49.1 (1.2)	43.4 (3.8)	71.6(1.2)

Table 2.6: F1 scores on test set, with SD in parentheses, for interpersonal emotion labelling task.

to civility is the only subtle linguistic cue that this tweet may constitute out-group speech.

- (3) a. Admire @OfficialCBC Chairman @Doe’s moral voice on issues of racism and restorative justice. He is a real leader for our nation and Congress.
- b. A decade has passed, but our friendship is the same. Proud to work with @Doe to #ReviveCivility. #tbt Read more about our efforts here:

Future work needs to look into what information the embeddings are using to make their classification decision.

Model Errors While the multitasking setup improves model performance on the task of predicting IGR, and outperforms human labelers in our small pilot, it still gets some easy examples wrong, such as labelling (4) as in-group even though it expresses some disgust at the target. The model also falls into the same trap as human labelers — for instance assuming that a tweet expressing admiration must be in-group (5).

- (4) Trump selected @USER for HHS Secretary. Price has undeniable history of cutting access to healthcare to millions, especially women.

Emotion	BERTweet	MultiTask
Admiration	77.9 (2.6)	72.8 (3.9)
Anger	71.7 (9.9)	69.4 (3.4)
Disgust	48.2 (22.4)	75.9 (6.5)*

Table 2.7: F1 scores on test set, SD in parentheses on out-group tweets. * indicates statistical significance ($p < 0.05$)

- (5) Inspiring speech from @USER - we have a duty to represent our country with respect & dignity. #NationalDayofCivility.

To ensure that model performance on IGR prediction is not limited by the size of our training data, we experimented with training BERTweet models on larger datasets. Since we have ‘found supervision’ for IGR labels, we only need to increase training data size by sampling more tweets from relevant accounts using the same procedure detailed in § 2.1.2. We found that F1 score does not increase with more training data.

Future work needs to look into linguistically motivated ways to improve model performance on the IGR task. Since we have observed that the multi-task setup improves model performance, perhaps other multi-task setups, such as modeling the overall affect towards the target expressed by the speaker might help in modeling IGR better.

2.4 DISCUSSION & CONCLUSION

Taking a cue from studies of bias in social science and psychology, we situate bias in language use through the lens of interpersonal relationships between the speaker and target of an utterance, and the speaker’s interpersonal emotional state with respect to the target. Over a corpus of tweets by members of US Congress, we introduce two novel tasks – intergroup relationship prediction (IGR) and interpersonal emotion labelling, to better understand variation in language as a function of social relationship between speaker and target in interpersonal utterances. We find certain interper-

sonal emotions like anger and disgust are over-represented in out-group situations, with the majority of the negative emotions directed at leaders of the two political parties. Through modeling studies, we find that transformer based models perform better than humans at predicting IGR given an utterance, raising the question as to what latent features of language the model uses to make this decision. Finally, we also find that joint modeling of the two dimensions is beneficial to prediction of certain interpersonal emotions in out-group situations.

This chapter lays the groundwork for future studies into intergroup bias more broadly. Analysis and modeling prove that intergroup bias can be systematically studied by decomposing it into discrete dimensions, and subtle variations can be captured using pre-trained language models fine-tuned for our task. Building upon the dataset and results in this chapter, we circle back to our motivations — **what are the systematic linguistic variations between in-group and out-group speech?**. The next chapter attempts to answer this question, through focused probing experiments into the neural models trained to discriminate in-group tweets from out-group tweets; Probing seeks to understand if the neural models learns to use human-interpretable higher-order linguistic features when performing the task.

Chapter 3: Intergroup bias gets a counterfactual probe

In the last chapter, we proposed a new framing of bias by modeling intergroup relationships (IGR, in-group and out-group) in interpersonal English language tweets, potentially capturing more subtle forms of bias. Analyses showed that this was a worthwhile framing of bias to understand it more comprehensively, but it raises a new question: which *linguistic features* vary systematically in different intergroup contexts?

As described in Chapter 1, LIB speculates that socially desirable in-group behaviors and socially undesirable out-group behaviors are encoded at a higher level of **abstraction**. The theory however relies on a restricted definition of abstractness that relies solely on predicates, and an ad-hoc analysis of ‘social desirability’ that doesn’t permit large-scale analysis. We can do better by using two well-defined pragmatic features: *specificity* (Li, 2017) is a pragmatic feature of text that measures the level of detail (similar to abstract–concrete axis), while *affect* is a feature that measures the attitude of a speaker towards their target (Sheng et al., 2019) in an utterance (analogous to social desirability).

Specificity and affect are analogous to the LIB axes of language variation that are easy to annotate and compute. Furthermore, specificity is a more *general* property than abstractness in the LIB — specificity is a property of the whole sentence rather than just the predicate. Thus, our study focuses intergroup bias more generally, rather than the narrow parameterization of the LIB. Similar to the LIB, our formulation of intergroup bias predicts that positive affect in-group utterances and negative affect out-group utterances are encoded with *lower specificity* (i.e. more generally). Tables 1.1 and 3.1 compare the predicted language variation between the LIB and our formulation.

In this work, we perform the first large-scale study of linguistic differences in inter-group bias by analyzing its nature in the corpus of English tweets described in § 2.1.2, which makes use of naturally occurring labels for in-group vs. out-group. This distinguishes us from existing work in LIB which mostly relies on artificial responses from participants in studies, rather than natural language use in the wild. To bolster our probing investigation, we also explore it causally: exploiting the quantitative nature of our formulation to study if a neural model fine-tuned for IGR prediction uses pragmatic features such as specificity and affect in its decision-making process through counterfactual probing techniques (Ravfogel et al., 2021).

To summarize our findings, we find a modest positive correlation between affect and IGR in our data, with a positive causation effect as well — making a tweet’s affect more positive makes it more likely to be in-group regardless of its specificity. We find no correlation between specificity and IGR in our data. Surprisingly, we discover a causal effect of low specificity on IGR prediction that is uniform across affect, but none for high specificity. We hypothesize that this could be because of damage to the underlying language model, but we leave further investigation to future work. We release our code and data [online](#).

3.1 PRIOR WORK

Intergroup bias The Linguistic Intergroup Bias (LIB) theory Maass et al. (1989); Maass (1999) tries to explain how stereotypes are transmitted and persist in communication by hypothesizing that socially desirable in-group behaviors and socially undesirable out-group behaviors are encoded at a higher level of abstraction. The LIB has been reproduced in various psychological experiments and analyses Anolli et al. (2006); Gorham (2006); it has also been used as an indicator for a speaker’s prejudicial attitudes Hippel et al. (1997), and racism Schnake and Ruscher (1998).

	in-group	out-group
positive affect	low specificity	high specificity
negative affect	high specificity	low specificity

Table 3.1: Predicted language variation in our more general formulation, using specificity and affect

Table 1.1 describes the LIB asymmetry and the parameters used. As stated earlier, the LIB relies on ad-hoc and hand-coded concepts such as ‘social desirability’ and abstractness of predicates (Semin and Fiedler, 1988). Our proposed experiments *generalize* beyond the LIB by utilizing parameters that are easily computable, and are a function of the whole utterance.

Specificity Specificity is a pragmatic concept of text that measures the level of detail and involvement of concepts, objects and events. Louis and Nenkova (2011) introduced the first dataset and model for sentence specificity prediction, and in later work Li (2017) illustrated the role of specificity in discourse coherence. Furthermore, Gao et al. (2019) expanded the scope of specificity analysis from the news domain to social media.

Affect There is a wealth of work studying emotions and sentiment in social media text (Mohammad, 2012; Wang et al., 2012; Mohammad and Kiritchenko, 2015; Abdul-Mageed and Ungar, 2017; Desai et al., 2020; Demszky et al., 2020). We introduced the first dataset annotated for *interpersonal* emotion (defined as only emotions expressed towards or in connection with a target), using the Plutchik wheel (Plutchik, 2001) as a framework. While fine-grained, this approach isn’t amenable to the experimentation we propose easily. Inspired by the concept of *regard* by a speaker towards a demographic in an utterance (Sheng et al., 2019), we introduce annotations for a

coarse-grained feature we term *affect* that estimates how a speaker feels towards the target they mentioned in an interpersonal utterance.

Table 3.1 describes the intergroup language variation as hypothesized in our experimentation, using specificity and affect. Analogous to LIB, our hypothesis is that positive affect utterances directed at in-group individuals, and negative affect utterances directed at out-group individuals are encoded with *lower specificity*.

AlterRep AlterRep (Ravfogel et al., 2021) is a probing technique that tests if a neural network *uses* a property, rather than just testing if the model’s learned representations correlate with the property. The method uses Iterative Nullspace Projection (INLP; Ravfogel et al., 2020) to iteratively train a linear classifier on the model’s internal representations to pick out a particular feature, using the parameters learned by the classifier to intervene on the embedding and alter it in a systematic way. The AlterRep method based on INLP has been used to probe for syntactic phenomena such as subject-verb number agreement (Ravfogel et al., 2021). To our knowledge, ours is the first work probing if a model learns and uses higher-level pragmatic features like affect and specificity using AlterRep.

3.2 PROBING FOR THE INTERGROUP BIAS

3.2.1 DATA & ANNOTATIONS

We use the same dataset of tweets from § 2.1.2, which consists of tweets by members of US Congress that @-mention other members in the same tweet, with ‘found-supervision’ for the IGR labels of every tweet. A tweet is in-group if it is targeted at another member of the same party as the writer of the tweet, else it is out-group.

Affect We build upon the dataset’s fine-grained annotations for interpersonal emotion by adding annotations for affect. We presented annotators on Mechanical Turk with tweets from our dataset with the target mention masked (with the placeholder Doe, to minimize potential biases of the annotator), and asked the following questions:

- a. How does the writer feel in general about Doe? *warmly, coldly, neutral, mixed*
- b. How does the writer feel in general about Doe’s actions/behavior? *approval, disapproval, neutral, mixed*

Annotators are given the option to select one of the 4 options listed above for each question. For each tweet, we collect annotations from 3 annotators, obtaining an aggregate label for each question by majority vote. We report an inter-annotator agreement score (Fleiss’s kappa; [Fleiss, 1971](#)) of 0.53 for the first question, and 0.56 for the second.

We derive a binary affect label (± 1) from our annotations using a simple rule: If the writer of a tweet is deemed to either feel warmly towards the target, or if they approve of the target’s actions, the affect is set to be positive; else it is set to be negative. An analysis of our collected annotations on the data shows that there is a small positive (Pearson’s) correlation ($r=0.2, p < 0.001$) between binary affect and IGR.

Specificity Specificity of the tweets in the dataset are calculated using the specificity prediction tool from [Gao et al. \(2019\)](#). Their specificity predictor is trained on tweets, and uses surface lexical features, as well as syntactic, semantic and distributional features to calculate a specificity score between 1 and 5. We note that on our dataset, there was *no correlation between specificity and IGR* ($r=-0.07, p < 0.001$), unlike affect. On further inspection of our dataset, we find that tweets with very high/low specificity scores (gathered by excluding specificity scores between 3 and

Affect	Train	Dev	Test
Positive	1813	226	242
Negative	589	80	83

Table 3.2: Distribution of affect in train-dev-test split

4, similar to excluding the middle in Gelman and Park, 2009) have a small but statistically significant negative correlation with IGR labels ($r=-0.13$, $p < 0.001$).

3.2.2 DATA & ANNOTATION

To obtain reliable annotations, we prequalify annotators using a qualifying task. Annotators were recruited on Mechanical Turk using a qualifying task where they were asked to annotate 6 tweets using the schema detailed in § 3.2.1. We restricted the qualification task to annotators living in the USA who had attempted at least 500 HITS and had a HIT approval rate $\geq 98\%$. After manual inspection, 6 anonymous annotators were qualified for bulk annotation. Each tweet was annotated by three different annotators. To ensure annotators were paid a fair wage of at least 10\$ an hour, we paid annotators \$0.50 per HIT. Each HIT involved annotating 3 tweets, which we estimate to take on average 3 minutes to complete. In total, 3,033 tweets between 2010 and 2021 were annotated with perceived affect.

We present preliminary statistics for the annotations on the dataset of tweets in Table 3.2.

3.2.3 INTERVENTIONS

Model We use BERTweet (Nguyen et al., 2020), a language model pre-trained on 850M English tweets, the same model used in § 2.3. All intervention experiments are carried out with the best performing *fine-tuned* version of this model — where the model is fine-tuned on the task of predicting IGR labels. The input to the model




Figure 3.1: Top-100 LM Accuracy on train plus validation split of our dataset for different intervention strengths.

is only the tweet with no other context, and the target masked with a placeholder @USER.

We use the model’s representations from layer 11 for the INLP procedure since it shows the most reliable effects. INLP (Ravfogel et al., 2020) works by learning a series of linear classifiers on the representations from an encoder. In each iteration, the embeddings are projected onto the intersection of nullspaces of the classifiers learned so far, meaning the information used by the existing classifiers is removed from the model. Every subsequent classifier we learn removes more information of the property of interest from the model’s representations. We find that higher layers offer a good balance between feature extractability and language model stability (see Figure 3.1) for our features.

After training INLP, AlterRep uses the classifier’s decision space to project model embeddings into a null component that contains no information from the feature of interest, and an orthogonal component, that contains all the information from the feature of interest. These two components thus enable us to perform the counterfactual intervention — pushing model embeddings towards having more, or less, of a particular property. When AlterRep uses INLP classifiers with more iterations, the strength of the intervention is greater. Figure 3.2 offers an illustration of our intervention experiment on specificity, and the expected results.

Affect Using the binary affect labels we derived from annotations that we described in § 3.2.1, we perform interventions to test if the model uses affect causally in its decision. We sample 3 tokens at random from each sentence in the training and validation split of our dataset, train an iterative linear classifier on the model’s representations of these tokens using INLP (against the affect label of the tweet), and use the decision boundary learned by the classifier to intervene by pushing model representations to have more positive affect or have more negative affect. We set the hyperparameter α in AlterRep to 4.

Specificity The INLP classifier for specificity is learned using the same procedure as for affect. We train the classifier on only the tweets with high and low specificity scores in our dataset (scores below 3 and above 4; scores taken from the specificity prediction tool in [Gao et al. \(2019\)](#)), excluding the middle to ensure effective learning of the decision boundary ([Gelman and Park, 2009](#)). Thus, we are effectively pushing the model representations to have high or low specificity. For both affect and specificity, once the INLP classifier is learned, we perform the intervention on a random subset of 30% of the tokens of a tweet (to control for tweet length). We also report the results of random interventions as a control, where random interventions are gen-

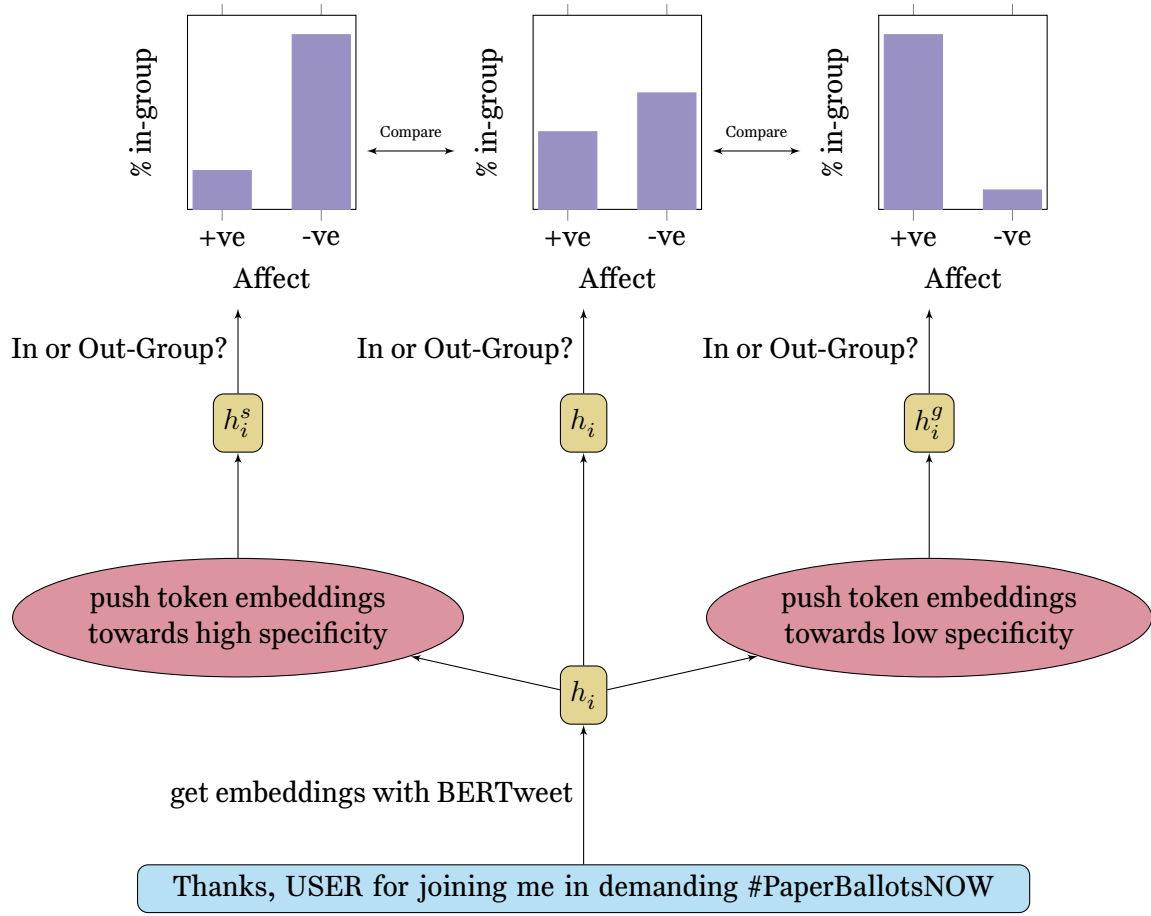


Figure 3.2: Flowchart describing the specificity intervention experiment and expected results.

erated by sampling from a standard gaussian instead of using the decision matrix generated by INLP.

Hypotheses We report the percentage of tweets in the test split of our dataset that are predicted to be in-group by our classifier model with increasing strength of the intervention (number of INLP iterations, 0 being pre-intervention). Thus, we have the following hypotheses on the effects of our intervention on the data based on our intergroup bias framework described in Table 3.1:

1. Interventions towards positive affect should induce the model to predict low specificity tweets to be in-group and high specificity tweets to be out-group, while interventions towards negative affect should affect the model conversely.
2. Interventions towards higher specificity should induce the model to predict positive affect tweets as out-group and negative affect tweets as in-group, while interventions towards lower specificity should affect the model conversely.

3.2.4 RESULTS & ANALYSIS

The results for the interventions on affect are presented in Figure 3.3, while those for specificity are presented in Figure 3.4. Overall, we observe that in both cases, interventions had the same effect on tweets that were annotated with positive affect as they did on tweets with negative affect (and similarly for tweets with high and low specificity) — so we only show the percentage of *all* tweets in the test split classified as in-group.

Affect As Figure 3.3 shows, pushing model representations towards having more positive affect causes almost all tweets in the test split of our data to be classified as in-group after 32 iterations of INLP. The randomness after 40 iterations of INLP could be attributed to the underlying RoBERTa language model being destroyed, as the LM Top-100 accuracy plot in Figure 3.1 shows. Pushing the model’s representations towards negative affect shows the inverse effect as expected, although the nature of the drop appears different. We hypothesize that this is because most of the tweets in our dataset (75.2%) have positive affect. An intervention pushing the representations towards negative affect would be slower and require stronger intervention forces, which is borne out in Figure 3.3.




Figure 3.3: % of test set classified as in-group plotted against number of INLP interventions for affect.

Specificity Figure 3.4 shows that pushing model representations to be more specific has no effect on model behavior and is indistinguishable from the control; but pushing towards lower specificity has a noticeable effect — interventions after 48 iterations of iNLP lead to all the data being predicted as in-group. Our hypothesis states that general language is more likely in positive affect in-group contexts; however we find no difference in the model’s behavior on positive versus negative affect tweets as reported earlier.

Overall our findings indicate that while the model does use affect towards making its decision on the intergroup relationship prediction task (albeit uniformly across specificity), it doesn’t use specificity as we had predicted. The discrepancy between high and low specificity interventions could be because the average specificity of tweets in our training data is 3.49 ($\sigma = 0.54$) — meaning that interventions towards lower




Figure 3.4: % of test set classified as in-group plotted against number of INLP interventions for specificity.

specificity act in opposition to most of our data in representation space. But these results requires further investigation to understand them better.

Qualitative error analysis Digging into the results further, we wanted to investigate if the interventions function the way we wanted them to. We analyzed the tokens that the model predicts before and after intervention for example (1). Firstly, fine-tuning the model for IGR prediction leads to degradation in LM abilities — a vanilla model predicts *birthday, anniversary* for the masked token in (1), but the fine-tuned model predicts nonsensical tokens like *sworn, opport__* even before any interventions. Pushing towards negative affect causes it to predict tokens with negative connotations (*killing, ass, opposition*), but degrades the underlying LM even further. The specificity interventions are especially hard to interpret due to the semantically and syntactically implausible tokens being selected (*opport__, mug__, ask__*)

(1) Happy <mask> @USER! I got you a new bill: #IIOA

While some of the interventions push the model’s predictions to be in the general lexical space desired (which probably explain the affect intervention results), the lack of contextual fit due to LM degradation may explain the inconclusive results, and lack of interaction between affect and specificity.

3.3 DISCUSSION & CONCLUSION

In this chapter, we performed a correlational and causal analysis of two pragmatic features, specificity and affect, on a dataset of interpersonal tweets in English, to establish how they influence intergroup relationship prediction. We find modest correlations between our features and IGR labels, while counterfactual probing reveals mixed results. Affect influences IGR prediction causally but without interacting with specificity, while specificity only influences IGR prediction in one direction.

Overall, we found no evidence for our general intergroup bias hypothesis, described in Table 3.1. While disappointing, the findings in this chapter are still instructive towards our next steps. There are 2 major issues that need to be addressed, if we are to truly discover systematic linguistic variations that differentiate in-group speech from out-group speech:

1. We need diverse real-world data, that reflects natural language use, at a larger scale. While our dataset of political tweets has served us well, it is limited in terms of its utility. Many of the tweets themselves are composed by social media teams employed by politicians, and are composed with political gain in mind. Moreover, the volume of tweets, and the representative groups in the context of U.S politics are small. We also need ample, high-quality data to build reliable computational models for prediction and analysis of the intergroup bias.

2. The data needs to be tied to real-world events. Currently, both axes of variation in our novel hypothesis, affect and specificity, are both features derived from the utterance. This confounds analysis and discovery of variation between groups — the LIB describes asymmetrical linguistic changes in abstractness in response to variation in acceptability of real-world behaviors (social desirability). We attempted to approximate social desirability through affect, but we don't have access to the direct affect of the speaker towards a referent, and this is impossible to obtain at scale. Thus, we need a robust measurement of a variable that describes the **state-of-the-world** preceding an utterance.

We tackle both of these issues directly in the next Chapter by introducing a novel dataset of online sports comments *parallel* with sports games — tying linguistic behavior to events directly.

Chapter 4: Intergroup bias in Reddit NFL game comments

*...a socially acceptable outlet for xenophobia. That is the function of organized sports
(in society) for the most part...*

—John Siracusa, [Reconcilable Differences Ep. 3](#)

Chapter 2 described my first data-driven study of intergroup bias in real-world language use, by curating a dataset of tweets by members of U.S. Congress with metadata derived gold labels for intergroup relationship. Crowd-sourced annotation revealed the systematic variation of interpersonal emotion with intergroup labels, and modeling further reinforced the systematic interaction between these two variables. In the previous chapter, I described my probing experiment to discover if the interaction between specificity and affect could describe any of the systematic *linguistic* variation observed in the data. While the results were inconclusive, they were instructive towards how we needed to proceed towards answering this question — we needed to look within utterances for intergroup variations in how people **referred** to the in-group and out-group, and we needed to *tie* the variation in intergroup language to the events preceding the utterance.

As I have described earlier, the LIB hypothesis is limiting and ad-hoc towards capturing the rich forms of variation that we do observe in natural language. For instance, the dataset of political tweets was restricted to tweets where there was only one explicit referent (via @-mentions). However, the **form of referencing** the in-group or out-group can reveal subtle biases as well. Consider these tweets:

- (1) a. **Mr. President-** Please tell your supporters to STAND DOWN, LEAVE the Capitol grounds and obey law enforcement who once again are risking their lives for our country!...

- b. ... Americans deserve answers on these unacceptable delays. We need a full accounting of **Pres Trump** and defense officials' decisions on Jan 6 ...
- c. We survived an insurrection and ... In we made it clear what St. Louis already knew: **Donald Trump** was the white supremacist-in-chief.

All of these tweets *refer* to the same individual — Donald Trump. However, (1-a) is by a Republican Congresswoman (and thus in-group), while (1-b) and (1-c) are by Democrats (out-group). One can observe distinct changes in how the speakers *refer* to him that is tinged by their (intergroup) relationships and changes in state-of-the-world. The Republican's tweet is more respectful, by only referring to him by position rather than name. The LIB hypothesizes variation only in the form of the predicate — this is an acceptable simplification when eliciting experiments in utterances, but discards useful information, as illustrated in these examples, in real-world language use.

To decipher the linguistic variables that underlie the systematic variation in the intergroup bias, we need to move beyond the political domain which suffers from two major drawbacks:

1. A difficulty in succinctly describing the events in U.S politics immediately preceding the tweet.
2. Much of the language by politician's is not natural speech by one person — it is the output of a social media team who monitor many if not all of the tweets.

In this chapter, I introduce a new dataset of interpersonal language — specifically sports comments from subreddits (internet forums) dedicated to fandoms for teams in the National Football League (NFL). As I will show, through careful data curation, we can obtain utterances with reliable information about the group allegiance of the writer of a comment (the in-group team they support), as well as a grounded score (the

win probability) that describes the state-of-the-world (in a non-linguistic manner) prior to the utterance of a comment. Annotation and preliminary analysis reveals a crucial blindspot in the LIB — the *form of the referent* (the argument to a predicate) that speakers use when referring to the in-group or the out-group may have systematic variations as well. By carefully exploiting the information processing capabilities of Large Language Models (LLMs), we can validate this systematic variation on a large-scale of over 200,000 Reddit comments spread over two years, revealing two striking social behaviors:

1. The better the state of affairs in the real world for the *in-group*, the more likely commenters are to **abstract away** from specifically referring to the in-group. This trend is remarkably linear across win probabilities for all types of in-group references.
2. References to out-groups by commenters remain stable over all win probabilities for the in-group, with only a slight uptick in the frequency of referring to the out-group using names or nicknames when they are close to defeat or victory.

These findings add much needed color to the LIB hypothesis — natural language is productive, and commenters can express their (implicit) intergroup bias beyond the predicate; The form of the referent itself shows systematic variation in response to changes in the state of affairs for the in-group. This chapter is an account of my efforts in this direction. § 4.1 gives an overview of the dataset, the structure of the NFL, and the robust nature of the grounded variables readily available against comments. In § 4.2, I detail the new protocol for tagging *references* to the in and out-group at a lexical level, to create a new expert annotated set of comments tagged with intergroup labels. Then, I move on to describing my efforts building models that can effectively learn to tag words/phrases with intergroup tags on a large-scale (§ 4.3), which enables the statistical analyses and insights I derive in § 4.4. I conclude in § 4.5 with a discussion of the findings in light of intergroup bias more generally.

4.1 A NEW DATASET OF INTERPERSONAL LANGUAGE

In Chapter 2, I listed two conditions for the type of language we want to analyze for intergroup bias, restated here:

1. Each utterance must have at least one target individual (person or group) about whom the utterance mainly concerns.
2. The relationship between the speaker and the target must be inferred based on metadata or other information.

These constitute **interpersonal** utterances. To these, we add another condition: a correspondence between each utterance and a **non-linguistic** description of the events that preceded (or precipitated) the events themselves. As I will show in this section, social media comments by fans of sports events, in particular the NFL, satisfy all of these conditions and offer a fertile ground for study.

4.1.1 PRIOR WORK

Language use within the domain of sports has been a rich source of analyses and studies within computational linguistics, including from the perspective of quantifying *social biases*. [Merullo et al. \(2019\)](#) studied commentator racial biases in descriptions of football players, reaffirming previous findings illustrating clear differences in terms of sentiment descriptions (white players were more likely to be described as intelligent), and name itself (white players were more likely to be referred to by their first name).

Human language learning and understanding does not happen in isolation; Indeed it is acquired and used in the physical world. Grounded language understanding aims to bridge the gap between the state-of-the-world, and the language that we use to

talk about it (Krishnamurthy and Kollar, 2013). The sports domain is suitable for exploring the link between language and grounded descriptions of the world, as sports like football employ scoreboards, statistics, and constantly updated databases to accurately track the state of the game. The state-of-the-world of a football game at any moment can be described using the score, the team in possession, yards gained, yards left to opponent touchline, and previous plays up to that point. Liang et al. (2009) rely on such descriptions to build a generative model that maps from utterances (in a recap of the game) to state-of-the-world (generated from the scoreboard and database of events provided by the NFL).

While there has been a wealth of work looking into the language used around sports and sports commentary, our work differs from previous work in two major ways. Firstly, ours is the first study to focus on the intergroup bias (rather than race or other social factors) — how do fans talk about their team, versus the opponent? Though sports as domain may seem trivial (as compared to politics, in our previous dataset), the insignificance of sports is precisely what makes studying human social-linguistic behavior in this realm interesting. People are less restrained speaking their mind freely, thus showcasing implicit (and explicit) prejudices freely ¹. Furthermore, **affective polarization is desirable** in sports for this very reason as well, whereas the rise of affective polarization has been studied extensively as a negative phenomenon in politics (Iyengar et al.).

Secondly, this dataset (and the analyses that follow) studies the intergroup bias grounded in the events of the game parallel with the online comments. Social desirability was chosen as one of the axis of variation in the original LIB — however, this is an ad-hoc formulation that is hard to generalize over and study at scale. Affect, as described in the previous chapter, was derived from the utterance itself rather than reflecting the state-of-the-world prior to the utterance. As I will explain, sports games, and in

¹as opposed to tweets by politicians (which are generally managed by social media teams) and commentators (who have to maintain an aura of neutrality, and obey broadcast regulations)

particular NFL football games, are rich with statistical information that allow us to describe the state-of-the-world on a numerical scale.

4.1.2 DATASET

Data source for intergroup comments Our new dataset of intergroup language comes from Reddit — specifically subreddits (forums) dedicated to fandoms for each of the 32 teams in the NFL. During the NFL season, each subreddit has *game threads* — posts created by moderators on which fans can comment in tandem with the game. Crucially, since every subreddit has their own thread, we effectively have a **parallel** intergroup language corpus; we have two teams and their fans, each with different allegiances, commenting on **the same events in the game**.

We focus on all completed games from 2021–22 and 2022–23 NFL seasons, and attempted scrape all comments from the game threads for both teams involved in every game. Furthermore, we also attempted to scrape comments from pre-game threads and post-game threads where available from subreddits. Table 4.1 gives summary statistics on our dataset after scraping. Within comments from game threads, we created a subset of comments that happened during active gametime — which we label as **gametime** comments. Most threads are closed, or become inactive, once a game ends. Game threads are usually created (by subreddit moderators) and open for comment a few hours before the start of a game.

Grounding football comments As remarked earlier, one reason for studying intergroup bias in sports comments is the ability to **ground** the language in quantifiable descriptions of the state-of-the-world. NFL, and American Football in particular, has some attractive features as a sport considering that our interest is in the *language* surrounding the events in the game. While physical, American Football is also one of the more strategic sports games, where outcomes are heavily dependent on a coach’s

Stat	Number
Teams	32
Games	568
Game threads	1104
Pre-game threads	261
Post-game threads	1040
Game thread comments	6,240,285
Gametime comments	6,679,988

Table 4.1: Summary statistics of our dataset. Game comments are judged to be game thread comments that happened *during* the course of the corresponding NFL game. Comment timestamps were compared with publicly available start and end times of games.

strategies and plays in a (relatively) small number of discrete events ([Pelechrinis and Papalexakis, 2016](#)).

There has been a wealth of work looking into predictive modeling of different statistics and events in a football game ([Horowitz et al., 2017](#); [Yurko et al., 2018](#)), from predicting expected points scored by teams, to yards gained by individual players. The state-of-the-world at any moment in a football game is determined by a variety of factors — the performance of teams before the game, the live score, the position of the offense, defense, and so many more. [Baldwin \(2021\)](#) modeled the Win Probability (WP) of a team at any point during the game using the following features:

- seconds remaining in half (and game).
- yard line
- score differential
- down, yards to go, timeouts remaining for each team

- whether team is playing at home
- Betting odds lines from Vegas

They find that simple decision tree based models with gradient boosting achieve a lower calibration error than previous models, and furthermore, incorporating the Vegas betting odds substantially reduced the error rate even further. For these reasons, we chose win probability (henceforth abbreviated as WP) as a succinct description of how desirable the state-of-the-world is to the in-group.

The NFL publicly releases play-by-play information after every game, which includes details on the plays and the timestamp of each play. The `nflFastR` package includes WP for the home team alongside each and every play, updating it as the game state evolves with each play. In concert with the comments (whose timestamp of submission we also have access to), we can thus derive the WP for the in-group at the time of the comment by selecting the WP for the most recent play at the time the comment was made. The WP for the away team is simply set as 1 minus the WP for the home team.

Table 4.2 lists some comments from the 2023 Super Bowl between the Chiefs and the Eagles, with the win probability for the Chiefs in the middle. As is evident, the WP cleverly models the complexities of a real-world sporting event into one number that accurately models how *desirable* the state-of-the-world is to the in-group. This a marked improvement on *social desirability* as an axis in the LIB, which was ad-hoc, and *affect* in our previous formulation, which was derived from the utterance rather than from the state-of-the-world.

4.2 TAGGING & ANNOTATION

The tweets in (1) illustrated a phenomenon that the Linguistic Intergroup Bias failed to account for — systematic variation in how people refer to the in-group versus the

Chiefs fans comments	Chiefs WP	Eagles fans comments
Now I'm nervous....	0.25	Good shit covey
Oh, is there a defense on the field?	0.75	Burn that clock baby

Table 4.2: Comments from the Chiefs and Eagles fans with the WP for the Chiefs in the middle. The WP for the Eagles is 1-WP for the Chiefs.

out-group. Moreover, our prior method of classifying the utterance(tweet) as a whole was too coarse to capture some of the variation observed in our original dataset, that we discarded due to the assumption of at most one -mention target. A preliminary analysis of our scraped data illustrate the problem with this approach:

- (2) a. Rams are gifting us a chance to win and we can't take advantage. The fuck!!!!
- b. if the ravens and chiefs beat these dudes by double digits then damn it so should we!

Even without contextual information for the above comments, we see a few different references to entities that we can readily identify as references to the in-group, out-group, and perhaps another category adjacent to the out-group. These examples suggest an alternate framing, or approach to modeling, of *references* to entities that are in-group or out-group, based on pre-existing tasks and pipelines in NLP: tagging.

4.2.1 TAGGING IN-GROUP VS. OUT-GROUP

Instead of judging an utterance as a whole to be primarily about the in-group or the out-group, we concern ourselves with how individuals are *referred* to in interpersonal comments themselves. The words or phrases that refer to the individuals can now be tagged with in-group([IN]) or out-group([OUT]) For instance, the examples in (2) could be **tagged** thus:

- (3) a. [OUT] are gifting [IN] a chance to win and [IN] can't take advantage. The fuck!!!!
- b. if [OTHER] and [OTHER] beat [OUT] by double digits then damn it so should [IN]!

In both sentences, we find that we can readily identify some words and phrases as references to the in-group with respect to the commenter — like the first person plural pronouns *we* and *us*. This is a common expression by sports fans to express affinity towards their in-group as highlighted in the language of these comments.

The spans ‘Rams’ in (2-a) and ‘these dudes’ in (2-b) are clear references to the *out-group* with respect to the commenter — one can come to these inferences with some reasoning over the choice of words by the commenter, verified further by knowledge of the source of the comments, and the live score. The spans ‘the ravens’ and ‘chiefs’ in (2-b) is more interesting — it is clearly not a reference to the in-group nor the opponent of the game. However, it is a reference to **a group of interest in this domain** — another NFL team and/or its fans. We consider these references to be [OTHER], and a special case of out-group references.

Tagging words and phrases within comments with their intergroup tags enables us to study utterance-level properties (how are people talking about the in-group and out-group), as well as how commenters choose to refer to the in-group and out-group themselves. However, to discover whether there are systematic differences in how commenters refer to and talk about groups, and the interaction therein with world-state(WP), we need a large, diverse sample of comments in our original dataset tagged. To build robust models for tagging comments with intergroup tags, we need a well annotated dataset of sufficient size and sample diversity. For this purpose, we construct a detailed annotation protocol with examples, which is described in the following section.

4.2.2 ANNOTATION PROTOCOL

Annotators are presented with a comment from our dataset, and some context from the state-of-the-world using a browser based interface. They are given the following high-level instructions:

1. All comments are from game threads corresponding to specific NFL games between two teams. You will be given the source of the comment — this is the team the writer of the comment supports, the opponent in that game, and the live score at the time of making the comment.
2. Highlight any words and phrases that refer to individuals (people, teams, subgroups within the team, organizations).
3. If the reference is to the same group as the source subreddit of the comment, tag this highlight as **in-group** ([IN]).
4. If the reference is towards the opponent in this specific game for which the comment is written, tag this highlight as **out-group** ([OUT]).
5. If the reference is towards any other team in the NFL apart from the two teams involved in this game, tag this highlight as **other** ([OTHER]).

The task of tagging words and phrases from comments in our dataset with intergroup tags can be highly involved, as the following examples show. In addition to knowledge of American Football, commonsense reasoning over the meaning of an utterance in context of the state-of-the-world (live score), one needs knowledge of the teams and players. For instance, in the following comment, one needs to know that the commenter supports the Seahawks, and that there is a prominent player named Wilson, to accurately tag in context that Wilson indeed is an in-group reference.

- (4) Our oline should start holding since apparently it 's okay now . Maybe Wilson can actually get some time to throw .

In other instances, the references to the in-group, out-group or other are not as explicit. However, we can infer based on context, and state-of-the-world (live score or WP), that the comment as whole, or a sentence in the comment, is **implicitly referring** to the in-group/out-group/other. Consider this example:

- (5) Lets go to the 4th with a 1st down around midfield.

There is no explicit word/phrasal reference to any team in the above comment². However, it is clear that the commenter is referring to the in-group — such an expression or admonition to the out-group or any other team would not be phrased in such a way. To facilitate these implicit annotations, we append a sentence-level token [SENT] before each sentence, and ask annotators to highlight and tag this sentence-level token if they believe the sentence as a whole is implicitly referential to a group. These annotations require a higher bar of reference, since all the comments are about the game at hand and will involve both teams to some extent. For instance, the following comments, we judge to not have explicit or implicit references to any relevant groups of interest (and label them as **null referent** comments henceforth):

- (6) a. Fair enough !
b. winning cures all lmao
c. turning the game off , have a good day yall

In case it is impossible to verify an explicit or implicit reference, annotators are instructed to not highlight any parts of the comment. While reasoning and information access (through team databases and search) can help in tagging several comments

²*lets* was originally a contraction of *let us* which has the first person plural pronoun.

Stat	Number
Game threads	768
Games	491
Comments	1499
Comments with no annotation	399
Number of <in> annotations	1393
Number of <out> annotations	266
Number of <other> annotations	166

Table 4.3: Summary statistics of expert annotated gold test set.

in the annotated dataset, pilot experiments revealed that a small fraction of comments are extremely hard to annotate, without onerous research into the events of the game and the live game thread. All annotation experiments were carried out using the `thresh.tools` annotation interface (Heineman et al., 2023).

Expert annotated gold dataset Due to the difficulty and involvement of this particular annotation task, we decided to rely on expert annotations for constructing a ‘gold’ annotated dataset. I personally annotated 1499 comments (randomly sampled from *game thread* comments) for intergroup references based on the protocol above. Some preliminary statistics of the annotated dataset are presented in Table 4.3. 26.7% of this random sample were judged to have no relevant intergroup reference, and in the remaining comments, references to the in-group vastly out-number references to the out-group or other groups. This is not surprising, since these are comments from forums dedicated to fandom of teams — people are much more likely to talk about their team over the opponent. This compliments our finding in Chapter 2 of in-group tweets being overwhelmingly positive.

4.2.2.1 Annotation Results & Analysis

To understand the annotation process better, and understand individual variation in what ‘counts’ as a reference to the in-group and out-group to be tagged, we recruited annotators for a crowd-sourced annotation experiment. This annotation experiment has two goals — gauge inter-annotator agreement among crowd annotators and sources of disagreement, and identify who operate under stricter constraints of time (and impoverished knowledge of the task and dataset)

Inter-annotator agreement To evaluate the reliability of collecting annotations for intergroup labels using the protocol above, we ran a small pilot study over a sample of 100 comments from the gold dataset. 3 annotators (undergraduates) were recruited to perform annotations, and presented with the same annotation protocol that I used to construct the gold annotation dataset. Average Fleiss K among these 3 annotators is 0.65, indicating moderate agreement.

In addition to the inter-annotator score, by counting exact-matches and weighting partial matches between individual crowd annotators and gold annotation, we calculate a score of 0.71 ± 0.03 . This gives a human upper bound for performance on this task, and characterizes its inherent subjectivity and difficulty.

Disagreement or diversity? Inter-annotator agreement among annotators, as well as the averaged accuracy of crowd-annotators against gold annotations, were both found to be in the moderate-to-high range of values. While this lends credence to the annotation protocol and the task design, looking at the source of disagreements can give us insights into the nature of the task itself, as well as why differences in judgements of intergroup affiliation can come down to annotator biases or judgement given context ([Atwell et al., 2022](#)), as well as implicit thresholding on what counts as a ‘reference’, as these examples demonstrate:

- (7) a. ...of the assorted ballparks and stadiums and concert venues I've been to
, **Lambeau** has the second worst bathrooms .
b. Can't do that against **an offense** this good.
c. Let 's go baby

Lambeau in (7-a) was judged to be a reference to the out-group in the context of this utterance — which was from the Colts subreddit in a game against the Green Bay Packers. However, to disambiguate this reference, annotators would need to know (or search to discover) that Lambeau Park is the stadium in which the Green Bay Packers play, and furthermore make a judgement call that this constitutes a relevant intergroup reference. As the expert annotator, I judged that it was a reference to the out-group because I discovered through search the relevance of *Lambeau* to the out-group.

In addition to requiring time and multiple steps of search and reasoning on the part of the annotator, the disagreements between annotators (and between annotators and gold annotation) illustrate variations in the implicit threshold of what constitutes an intergroup reference. For instance, *an offense* in (7-b) was judged by some annotators to refer to the out-group in context (the in-group was losing), however the generic nature of the referent lead some annotators to judge that this was an overall statement about the game, rather than a reference in of itself. Similarly with (7-c), some annotators judged the comment as a whole to be about the in-group (boosting or building them up) based on context and commonsense reasoning; However, other annotators judged this to be an expression of excitement by the fan, and not an obvious reference.

Is disagreement a signal of intergroup bias? Overall, these findings and examples paint a complex picture of how to reconcile agreement between annotators, and disagreements that nevertheless provide reliable signal towards the intergroup bias. Whether or not examples in (7) contain references to the in/out-group is not simply a conse-

quence of the difficulty of our task, or the inability for annotators to transparently describe the mental state of commenters (which we also observed annotating for interpersonal emotion as described in § 2.1.3). Rather, we need to analyze them as possibly another subtle influence of the intergroup bias itself — demonstrated by posing the question as to why annotators chose the forms in (7) rather than in (8), which seemingly convey the same meaning, and would be uncontroversial in annotation:

- (8) a. ...of the assorted ballparks and stadiums and concert venues I’ve been to
 , **the Packer’s** stadium has the second worst bathrooms .
 b. Can’t do that against **the Packers** offense this good.
 c. Let’s go **Packers**!

The intergroup bias can be observed to lead to subtle shifts in reference form, as we will see in forthcoming sections. However, future work needs to look into embracing the inherent annotation diversity in this task with grounded context, which can boost model performance on involved tasks that require linguistic and commonsense reasoning over world knowledge and context ([Atwell et al., 2022](#)).

4.2.3 QUALITATIVE ANALYSIS & TRENDS

The annotated dataset enables us to study qualitative trends, that will guide quantitative modeling and regression studies presented in § 4.4. I want to specifically focus on two phenomenon that are directly observable in the data and illustrated with examples — diversity in form of referring expression, and trends over WP.

Mereology of referring expressions Expert annotation revealed that commenters refer to groups of interest in a myriad of different ways. In the previous section, we liberally defined the annotation protocol for highlighting references to *individuals* in




Figure 4.1: Frequency of any-group, in-group and null references in comments that fall in 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show some noisy trends.

the in-group, out-group and other. Using insights from mereology (Varzi, 2019), I derive a taxonomy of ‘parthood’ in intergroup relations, that defines what it means for a reference to constitute a reference towards the in-group/out-group/other:

1. **Names of people:** Commenters frequently refer to individual players and coaches using names, nicknames, shirt numbers, initials, pronouns, etc. : *Tua*, *TK87*, *he/him*...
2. **Subset of the team:** This refers to groups of players, or coaches, rather than just one player: *the offense*, *our defense*, *o-line*, ...
3. **Team:** This is the standard way of referring to the team, but there is a host of variety within this category as well. In addition to the name of the team (*rams*,

bills, cowboys), nicknames (*lambs, cowgirls*), city names (*LA, Buffalo, Dallas*), commenters also use pronominal expressions like *our boys* for the in-group, pronouns like *they/them* for the in-group and out-group, and many more.

4. **Team plus supporters:** This is the largest possible reference to the in-group/out-group. References to the in-group generally involve the first person pronouns *we* and *us*, but can also be done with the third person pronouns *they* and *them*. The latter of course, could also refer to out-group or other, and require contexts to disambiguate.

The taxonomy above is ordered ad-hoc in order of increasing coverage of the whole group, by the referential part — the size of the reference gets larger from people to the entire group. Thus, players are the smallest unit of reference within a group, and the team/organization plus its supporters constitute the largest possible reference to the group itself. However, there are diverse ways of referring to each group within each, as the examples above illustrate.

Trends Within this expertly annotated dataset, we can observe two clear trends by plotting the frequency of a feature of interest over comments that fall within a win probability (WP) window. Figure 4.1 plots the frequency of any reference (in-group, out-group or other), in-group references, and null references over all WP probabilities:

1. References to the in-group, and references to any group overall, clearly go down with WP.
2. Null references increase steadily with WP.

The steady increase in number of null references in higher WP windows is interesting, and seem to be attributable to a combination of excitement at being close to victory,

and a tendency to abstract away from specific events to be terse and celebrate the victory, as the following examples show.

- (9) a. I can’t stop smiling!
- b. Absolute domination!
- c. ITS COMING BACK!

While the trends observed in this section are not statistically or practically significant (the slopes in Figure 4.1 are small), this can be attributed to the small sample size of analysis. Our expert annotated dataset is only 1500 comments. The intergroup bias is a social phenomenon, and like many social phenomenon, we can make clear inferences at scale. Obtaining human annotated data at scale would be prohibitively hard and expensive in this setting — thus, we turn to fine-tuning Large Language Models, to automate this task, thus allowing us make inferences about trends in the intergroup bias as a function of WP. The next section will also give us insights into the difficulty of training large machine learning models on the difficult task of *tagging* words and phrases in comments with their intergroup referential tags.

4.3 MODELING THE INTERGROUP BIAS WITH LLMS

Large Language Models (LLMs) have shown remarkable abilities in various domains over the last few years (Brown et al., 2020). LLMs trained on next-token prediction score highest on numerous benchmarks: from linguistic understanding (Srivastava et al., 2023), knowledge and search (Team, 2023), to complex reasoning (Wei et al., 2022). Recent very large language models exemplify *in-context learning* (ICL) behaviors as well. Rather than finetuning on specific tasks, ICL relies a few training examples (few-shot) fed into the model prompt at inference time, effectively “learning” via inducing internal activations of the LLMs in such a manner as to pre-dispose it to perform well at the task.

As I’ve shown in the previous section, annotating comments from our dataset to highlight spans that refer to the in-group, out-group or other requires linguistic understanding, knowledge of the NFL and its teams, as well as complex reasoning over why a commenter might choose certain word forms compatible with the state-of-the-world — LLMs excel at all of these. Can we use LLMs to automate tagging of comments in our dataset with intergroup references? Consider the following example, with information about the in-group w.r.t the commenter, the out-group (opponent), and the win probability at the time of making the comment:

COMMENT: If we could combine Pickett 's underthrows with Mariota 's overthrows , we'd get a pretty good QB.

IN-GROUP: Falcons

OUT-GROUP: Steelers

WP: 0.12

The first person plural pronoun *we* should be annotated as in-group — a model trained to tag comments on our dataset must learn that fans use *we* to signal their group membership and fandom. There are two proper names denoting players; A model resolve the last names to specific players, figure out the team they play for, and thus tag the names respectively with the correct tags. The model could also reason over the structure of the sentence and the aligned WP, and figure out that Pickett is probably a member of the out-group, while Mariota is probably in-group, since the commenter wants to pick and choose qualities from another player to help with their own team’s performance while they are losing. The final output should be of this form:

TAGGED COMMENT: If [IN] could combine [OUT] 's underthrows with [IN] 's overthrows , [IN]'d get a pretty good QB.

This is a complex task, and for this reason we design a modeling approach that aims to exploit the capabilities of LLMs to the fullest extent possible, and understand

which techniques serve us best towards tagging on a large scale. We focus on fine-tuning an off-the-shelf Instruction-tuned encoder-decoder language model Flan-T5-Large (Chung et al., 2022) specifically for our task, and analyze the impact of instructions, few-shot examples, and explanations on the task of tagging comments with intergroup tags.

4.3.1 MODEL TYPES

Flan-T5-Large is an encoder-decoder model — the encoder encodes the input, upon which the decoder learns to attend to different regions of interest when generating its output, token by token. Thus, our task is framed end-to-end as the model receiving as input the untagged comment with contextual information (in-group, out-group, WP), and being asked to generate the comment with relevant words/phrases replaced with the appropriate tags. For instance, here is a sample training input:

```
COMMENT: [SENT] Defense getting absolutely bullied by  
a dude that looks like he sells solar panels  
IN-GROUP: Jets  
OUT-GROUP: Bears  
WIN PROBABILITY: 71.5%  
TARGET:
```

and here is the model’s expected output after the word TARGET:

```
[SENT] [IN] getting absolutely bullied by [OUT] that  
looks like [OUT] sells solar panels .  
REF_EXPRESSIONS: ['Defense', 'a dude', 'he']
```

In addition to the tagged comment, the model is trained to generate a list of the referring expressions that have been tagged in the original comment (for ease of analysis later).

In general, performance of LLMs has scaled with compute and data in recent years (Kaplan et al., 2020). However, the best LLMs are prohibitive to use in academic research settings due to cost and resources. Finetuning Flan-T5 by itself solely on our gold dataset with the above format leads to poor performance by itself, since our training dataset is small. However, we can bolster its performance by incorporating a few different techniques:

Few-shot To facilitate ICL, we prepend 5 examples (not from our gold dataset) with the above training datapoint format to each datapoint in our training data. This gives implicit demonstrations to the model of the behaviors we want it to learn.

Instructions We prepend a short paragraph of text, giving instructions on how to perform the task, similar to the protocol we provided annotators. This is to further guide the model’s ICL activations towards learning the tagging behavior correctly.

Explanations Chain-of-Thought (CoT) prompting (Wei et al., 2022) has been shown to elicit complex reasoning for LLMs, by giving the model’s a larger scratch-pad and guiding model activations further towards desired behaviors. While CoT has generally been observed in extremely large models beyond the scope of this study, Wadhwa et al. (2023) show that finetuning with CoT explanations generated from a larger model improves performance on relation extraction. We append a small explanation to the few-shot examples that explains why the comment ought to be tagged in this manner, and generate explanations for all datapoints in our gold dataset. The model is thus tasked to generate an EXPLANATION after REF_EXPRESSIONS.

To generate explanations for all datapoints in the dataset, we use `gpt-3.5-turbo`, which we prompt with instructions and few-shot examples with explanations appended. Here is an example of a GPT generated explanation for the training example at the beginning of this section:

The commenter is likely a fan of the Falcons, as they are playing against the Steelers and have a low win probability. The comment is criticizing the quarterback play of both teams, mentioning 'Pickett' (Steelers) for underthrows (tagged as [OUT]) and 'Mariota' (Falcons) for overthrows (tagged as [IN]). 'we'd' should also be tagged as [IN] since it refers to the in-group, the Falcons.

GPT-generated explanations involve complex reasoning over the entities mentioned in the comment, knowledge of football, as well as linguistic reasoning. Thus, while our model is relatively small ³, we can still gain the advantages of CoT through fine-tuning on GPT-3.5 generated explanations.

Ceiling performance To compare the effectiveness of our techniques, we also evaluate the performance of GPT-3.5 and GPT4 ([Achiam et al., 2023](#)) on our gold dataset. We present each datapoint from our dataset with instructions, few-shot examples and explanations, asking the model to generate an explanation, a list of referring expressions, and the target tagged comment in that order.

Evaluation & Implementation Based on the techniques defined above, we build and compare 4 different fine-tuned models that explore 4 different combinations of tech-

³768M parameters, compared to GPT-3.5's 175B parameters

niques: fewshot, fewshot plus instructions, fewshot plus CoT explanations, and few-shot plus instructions plus CoT explanations. We partition our gold dataset into a test set of 318 datapoints, and a training set of 1181 datapoints.

To evaluate the performance of a model on the test dataset, we employ two forms of metrics:

- We calculate **micro-F1** scores for each of the tags across the whole test set. Sometimes the model accurately tags the correct word/phrase, but is off by 1-2 characters/words. For this reason, we weigh each true positive match: if it is within one character of the gold tag’s location, we assign a full correctness score of 1, if it is within 3 characters, a score of 0.5, and if it is within 5 characters, a correctness score of 0.25.
- We calculate two automated metrics that compare two sequences and generates the — GLEU score (Napoles et al., 2015) and word error rate (WER) (Woodard and Nelson, 1982). The former was designed for evaluating grammatical error correction, while the latter is analogous to edit distance. Both metrics are suitable for our task, where we want to penalize large amounts of differences between model generated tagged comments and the gold comments, while copying over most of the input comment.

4.3.2 RESULTS

Table 4.4 shows the overall results on all fine-tuned models as the GPT models. We observe that the model employing a combination of all of our techniques performs best overall, beating GPT-3.5, a model 2 orders of magnitude bigger, and is on-par with GPT-4⁴. The automated metrics reveal that each additional technique on top of

⁴rumored to be a Mixture of 8 Experts each 175B in size

Model	F1	WER	GLEU
fewshot	56.9	5.5	88.5
fewshot+instruct	54.0	5.2	88.1
fewshot+cot	54.5	5.3	88.1
fewshot+instruct+cot	60.0	4.9	89.8
gpt-3.5-turbo	48.4	16.2	72.6
gpt-4	60.6	7.4	86.4

Table 4.4: micro-F1 score (higher is better) over all tags, WER (lower is better) and GLEU scores (higher is better) on test split for all models.

few-shot examples reduces error ever so slightly, with the fewshot+instruct+cot model performing best in both automated metrics and tagging accuracy.

Copy errors The low performance of the GPT models on the automated metrics is a function of low performance on copying behavior. GPT-3.5 and 4 need to copy most of the input comment over with small changes only over words and phrases to be tagged. Due to the non-deterministic and stochastic nature of these models, they make small errors in copying behavior — missing some words, adding new words, changing tense/aspect of other words. Fine-tuning leads to much lower errors in copying.

Low recall Table 4.5 presents the per-class recall performance of all models, illustrating the shortcomings of the fine-tuned models as compared to the much larger GPT-3.5 and GPT-4. GPT-3.5 and 4 perform better than our fine-tuned model on recall primarily for **tagging out-group references**, where our fine-tuned models overgeneralize towards tagging proper names as in-group (or not tagging it at all) due to the tag imbalance in our training data. GPT-4 correctly tagged the bolded names in

Model	[IN]	[OUT]	[OTHER]
fewshot+instruct+cot	66.6	32.1	14.1
gpt-3.5-turbo	54.5	39.3	32.1
gpt-4	62.5	53.6	28.8

Table 4.5: Recall scores for each tag on test split for all models.

the following examples as out-group, with the generated explanation also correctly citing

- (10) a. The fact that we dont have 10 sacks is just a testament to **Josh Allen**.
- b. Anyone have a clip of the hit on **Huntley** ?

Is the model’s performance affected by win probabilities? To ensure the analyses that follow are not an artifact of the model’s weakness on comments from certain win probabilities, we verified that there was no correlation between the model’s performance and win probability.

4.4 LARGE-SCALE ANALYSIS OF MODEL TAGGED COMMENTS

In § 4.2.3, we observed a clear intergroup bias over state-of-the-world (WP), however the significance of these findings were hampered by the same size of our hand-annotated data, and the expense of scaling towards tagging more comments. Our best fine-tuned model performs extremely well at tagging references to the in-group; while its performance on out-group (and other) references is low, the low-recall is uniform across win probabilities. Thus, the analyses presented in this section are still demonstrative of subtle manifestations of the intergroup bias on real-world language use, that could previously not be described by the LIB.

For the following analyses, we sampled over 223,680 comments from game threads, specifically focusing on **gametime comments**. We used the best performing fine-tuned




Figure 4.2: Frequency of any-group and null references over all 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show clear linear trends.

model from § 4.3.2 to generate tagged comments, as well as the list of referring expressions and their corresponding intergroup tags. To mitigate some errors in copying and displacement, we performed edit checks between the tagged comment and the original comment so that the list of referring expressions was valid. In addition, we performed heuristic based tagging on top of the tagged comments from the model for first-person plural references using *we/us*, as well as common names and nicknames for teams.

In-group trends Figure 4.2 plots the frequency of occurrence of references to any group(*any_ref*) and null references (*none*) over all WP from 0 to 100. Within each win probability window of 5 percentage points, we count the occurrence of a vari-




Figure 4.3: Normalized frequency of references to the in-group, in-group using *we/us* and out-group over all 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show clear linear trends.

able of interest, and normalize it by the number of comments within that window. As the figure shows, there is a steady decline in the frequency of references overall, and a steady increase in the frequency of null references. The trend is surprisingly linear in its behavior, with outliers at the lowest, and highest win probability windows, where deviations in linguistic behavior by commenters might be expected, due to the certainty of winning and losing respectively.

A cursory analysis of high WP(see (11)) and low WP(see (12)) comments with no references reveals an obvious increase in positive sentiment, but also increased terseness closer to victory. Comparing the length of comments, we find a very small but significant ($\rho = -0.05$ ($p < 0.005$)) negative correlation between length of comment and the win probability.




Figure 4.4: Normalized frequency of references to the in-group and out-group solely using names or *they/them* over all 5% WP windows from 0 to 100. A simple regression line with 95% CI is fit separately for each feature to show clear linear trends.

- (11) a. HOLY SHIT
b. DO NOT TAKE YOUR FOOT OFF THE GAS
c. WHAT A THROW
- (12) a. Lol like a preseason game.
b. Yeah I 'm done for tonight
c. Bruh ...

Overall, this seems concurrent with the idea that the better the state-of-the-world is for the in-group, the less likely commenters are to refer to the in-group in any form. Commenters are more likely to abstract away from referring to the in-group (or out-

group) using explicit or implicit mentions, towards general expressions of sentiment or descriptions of the game itself (like in (11)).

The in-group sticks together in solidarity Figure 4.4 plots the frequency of references to the in-group and out-group, normalized by the number of comments that contain any reference at all — this is done to understand variations between in-group and out-group references as reference frequency goes down overall. Even controlling for the overall decline, we notice a marginal decline in (normalized) in-group reference frequency, and a steady increase in (normalized) out-group reference frequency. The more likely the in-group is to win, references to the in-group go down while references to the out-group go up.

- (13) a. If **we** get a stop here and a touchdown on the next drive it 's a ballgame
let 's fucking go .
b. wish **we** could fire fangio this week...

Figure 4.4 also plots references to the in-group using first-person plural pronouns (labelled *we*), while Figure ?? plots relative frequency of reference to the in-group or out-group using names (*in-names* and *out-names*) or third-person plural pronouns (*they_in* and *they_out*). Overall, these paint a complex picture. Pronouns like *we* and *they* carry a lot of affect, and are associated with inclusivity and othering. Commenters are more likely to use *we* in a losing scenario, and use it more often than in a winning scenario to build solidarity and hope, as well as offer criticism, as examples in (13) show. As WP increases, references to the out-group using all forms goes up

Feature	Slope	r-squared
Any reference	-17×10^{-4}	0.61
No reference	3.5×10^{-4}	0.57
In-group	-2.6×10^{-4}	0.31
we	-4×10^{-4}	0.87
out-group	2×10^{-4}	0.33
in-group names only	-2×10^{-5}	0.04
out-group names only	1×10^{-4}	0.58
they_in	-4×10^{-5}	0.22
they_out	6×10^{-5}	0.33

Table 4.6: Table of slopes of feature of interest against increasing WP, alongside the r-squared showing how much of the variance is explained by the linear regression fit.

4.5 DISCUSSION & CONCLUSION

There are two takeaways from the application of our best fine-tuned model on a large, representative sample of our dataset. First, is the **linear and inverse** relationship between the frequency of references to the in-group and the state-of-the-world (win probability — WP). Table 4.6 showcases the slope and percentage of variation explained by a linear fit over only WP, reinforcing this finding. Second, and concomitant with the previous finding, is the relatively stable association between frequency of out-group references and WP, and the **marked increase** in non-referential comments with increases in WP. We shall discuss both of these findings and their implications towards the intergroup bias in turn.

In-group Figures 4.2 and 4.4 display a remarkable linear relationship between two parameters that are operationalizations of the different aspects of the world. What does the decrease mean in context of this domain, and what is the significance of its linearity? This finding adds further evidence that not only is WP an accurate description of a complex state of affairs in football games, it is sufficient to capture a large amount of variation in linguistic behavior. In addition to being well calibrated as a

machine learning model as described in [Baldwin \(2021\)](#), these results show that it is also well calibrated to **human appraisals of the state-of-the-world**. Through well devised ‘scoreboards’ based on non-linguistic variables of interest, this gives us hope that similar simplified scores can be calculated to describe the state-of-the-world in non-sporting domains, towards replicating these findings in conversation more generally.

The trends observed with in-group references versus WP also add to the subtle ways we perpetuate bias in our linguistic behavior, especially in this case towards **in-group preservation** ([Maass, 1999](#)). While commenters are more than willing to criticize the in-group across WP, the self-protective instinct is evident in the way they choose to refer to the in-group using *we/us* rather than *they/them*, or to not refer to the in-group at all, abstracting away to talking about their sentiment or description of the events. How commenters choose the form of reference to an in-group constitutes just as subtle a bias as their choice of predicate.

Out-group Regarding the out-group (and other), normalized reference frequency increases steadily — however commenters also refer to the out-group explicitly or implicitly much less. This is similar to our findings in Chapter 2, where there was a clear bias in positive emotions towards the in-group, revealing a drawback of studying intergroup bias on naturally occurring language (as opposed to elicited utterances generally in the LIB) at scale — speakers prefer to talk about their in-group overall, and with negative affect when they do talk about the out-group. References to the out-group are most frequent at the lowest/highest WP (corresponding to the end of a game when win/loss is certain) with a linear increase in-between.

Modeling Improvements As discussed in the previous section, there is room for improvement on large-scale tagging of referential expressions, especially out-group references. This doesn’t detract from our analysis in the previous section — the model’s

performance isn't correlated with WP, and a significant reason for the low performance of our model can be attributed to size. GPT-4 outperforms our model on recall over out-group references, and a lot of its gains come courtesy of its size. A new generation of open LLMs including Mistral ([Jiang et al., 2023](#)), Gemma Team ([2023](#)), and OLMo ([Groeneveld et al., 2024](#)) provide competitive performance on benchmarks to GPT-4, and exhibit complex reasoning and general knowledge of events and entities — crucial towards performance on our tagging task. These need to be explored for further gains on out-group tagging performance. Furthermore, incorporating annotator confidence into modeling directly can boost performance ([Atwell et al., 2022](#)), as well as help the model learn reliable signals from intergroup bias as we observed in § 4.2.2.1.

Chapter 5: Summary

This dissertation is a compilation of three different data-driven experimental studies into intergroup bias in natural language online.

Intergroup bias in tweets Chapter 2 introduced the idea of studying intergroup bias more broadly as distinct from the artificial conditions under which the LIB has proven to exist. By decomposing the bias into two constituent dimensions — emotion and intergroup relationship, we can observe systematic relationships between these two. Modeling further fortifies this correlation, showing that one can derive insights into intergroup bias through modeling and analysis of natural language use online.

Counterfactual probing for intergroup bias Building atop the dataset and models in Chapter 2, in Chapter 3 we investigated if the neural models fine-tuned for predicting intergroup relationships use human-interpretable higher order semantic features like affect and specificity in making their decisions. Counterfactual probing techniques allowed us to ask this question in a measured way on neural models. While our specific hypothesis failed to hold up, our experiments revealed the need for grounded language data in order to tie systematic variation in in-group and out-group speech to events that preceded them.

Grounding intergroup bias in football comments Chapter 4 introduces a new dataset of interpersonal comments from live game threads on NFL team subreddits during Football games. By cross-comparison with historical statistics from live games, we can use win probability as an accurate measurement of state-of-the-world leading up to comments by participants about the in-group or out-group. We find several significant, linear patterns in the form of referent used to refer to in-group and out-group;

The more likely the in-group is to win, the more likely fans are to abstract away from referring to the in-group by name, pronoun or even implicitly, towards abstract descriptions of the events with no clear referent. Referents to the out-group however, remain stable across win probabilities.

Overall, the findings in this dissertation refine the LIB for broad application and study in real-world language use, by encompassing more forms of variation in the entirety of an utterance, and including other axes of human expression, like affect and emotion.

Implications for contemporary work on bias

5.1 FUTURE WORK

This dissertation is an attempt at answering complex questions about subtle changes in linguistic behavior. There are several unanswered questions, and promising topics for future investigation, a few of which I summarize below.

Stereotypes The LIB hypothesis is a mechanism for how bias subconsciously seeps into utterances; However, the impetus for the hypothesis was to explain how stereotypes persist and propagate in communication. This was the reason for the focus on *abstractness* of predicates — there is some evidence that ‘abstract’ language (as narrowly defined) is liable to reinforce conventions and labels cognitively. By relaxing the focus on stereotype-related behaviors and speech, we have been able to study more natural language; But this renews the question of what effect the observed linguistic behaviors have on stereotype perpetuation.

Parallel & Multilingual Chapter 4 introduces a dataset of interpersonal comments, with parallel live-game statistics to ground the utterances. Beyond the findings in

§ 4.4, the dataset has a wealth of information in the *parallel* nature of the utterances from fans of both teams observing the same sequence of events. Furthermore, there are possibilities for expanding the scope of this study to multilingual parallel comments, for instance from the live-broadcast comments of NFL games in Spanish and English.

Language Modeling with grounding As the findings in Chapter 4 reveal, win probability (WP) is a remarkable real-valued indicator of the state-of-the-world, compressed succinctly in a real-valued number. Crucially, it grounds the utterance for both sets of fans and their in-group robustly, evidenced by the linear relationship between WP and several linguistic indicators. While a wealth of work has looked into grounding language models in images and other modalities ([Radford et al., 2021](#)), there has been relatively little work looking into grounding with parallel meaning representations of ground state. The parallel corpus of game comments from two perspectives, with a non-linguistic description of events, should prove to be an exciting avenue for research into training language models with steerable intergroup dynamics, or perspective.

References

- Muhammad Abdul-Mageed and Lyle Ungar. EmoNet: Fine-Grained Emotion Detection with Gated Recurrent Neural Networks. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 718–728, Vancouver, Canada, July 2017. Association for Computational Linguistics.
- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Roei Aharoni and Yoav Goldberg. Unsupervised Domain Clusters in Pretrained Language Models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7747–7763, Online, July 2020. Association for Computational Linguistics.
- Luigi Anolli, Valentino Zurloni, and Giuseppe Riva. Linguistic Intergroup Bias in Political Communication. *The Journal of General Psychology*, 133:237 – 255, 2006.
- Katherine Atwell, Remi Choi, Junyi Jessy Li, and Malihe Alikhani. The role of context and uncertainty in shallow discourse parsing. In Nicoletta Calzolari, Chu-Ren Huang, and et al., editors, *Proceedings of the 29th International Conference on Computational Linguistics*, pages 797–811, Gyeongju, Republic of Korea, October 2022. International Committee on Computational Linguistics.
- Ben Baldwin. Open source football: nflfastr ep, wp, cp xyac, and xpass models, 2021.

Francesco Barbieri, Jose Camacho-Collados, Luis Espinosa Anke, and Leonardo Neves. TweetEval: Unified Benchmark and Comparative Evaluation for Tweet Classification. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 1644–1650, Online, November 2020. Association for Computational Linguistics.

David Beaver and Jason Stanley. Toward a Non-Ideal Philosophy of Language. *Graduate Faculty Philosophy Journal*, 39(2):503–547, 2018.

Andrea Beltrama. Social meaning in semantics and pragmatics. *Language and Linguistics Compass*, 14(9):e12398, 2020.

Taylor Berg-Kirkpatrick, David Burkett, and Dan Klein. An empirical investigation of statistical significance in NLP. In *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, pages 995–1005, Jeju Island, Korea, July 2012. Association for Computational Linguistics.

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeff Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. *ArXiv*, abs/2005.14165, 2020.

Hyung Won Chung, Le Hou, S. Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Webson, Shixiang Shane Gu, Zhuyun Dai, Mirac Suzgun, Xinyun Chen, Aakanksha Chowdhery, Dasha Valter, Sharan Narang, Gaurav Mishra, Adams Wei Yu, Vincent Zhao, Yanping Huang, Andrew M. Dai, Hongkun Yu, Slav Petrov,

Ed Huai hsin Chi, Jeff Dean, Jacob Devlin, Adam Roberts, Denny Zhou, Quoc V. Le, and Jason Wei. Scaling instruction-fine-tuned language models. *ArXiv*, abs/2210.11416, 2022.

Dorottya Demszky, Dana Movshovitz-Attias, Jeongwoo Ko, Alan Cowen, Gaurav Nemade, and Sujith Ravi. GoEmotions: A Dataset of Fine-Grained Emotions. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4040–4054, Online, July 2020. Association for Computational Linguistics.

Shrey Desai, Cornelia Caragea, and Junyi Jessy Li. Detecting Perceived Emotions in Hurricane Disasters. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5290–5305, Online, July 2020. Association for Computational Linguistics.

Penelope Eckert. Three waves of variation study: The emergence of meaning in the study of sociolinguistic variation. *Annual review of Anthropology*, 41: 87–100, 2012.

Joseph L Fleiss. Measuring nominal scale agreement among many raters. *Psychological bulletin*, 76(5):378, 1971.

Yifan Gao, Yang Zhong, Daniel Preoțiuc-Pietro, and Junyi Jessy Li. Predicting and analyzing language specificity in social media posts. *Proceedings of the AAAI Conference on Artificial Intelligence*, 33(1):6415–6422, 2019. ISSN 2374-3468. Number: 01.

Andrew Gelman and David K. Park. Splitting a predictor at the upper quarter or third and the lower quarter or third. *The American Statistician*, 63(1):1–8, 2009.

Bradley W. Gorham. News Media’s Relationship With Stereotyping: The Linguistic Intergroup Bias in Response to Crime News. *Journal of Communication*, 56(2):289–308, 2006. ISSN 1460-2466(Electronic),0021-9916(Print). Place: United Kingdom Publisher: Blackwell Publishing.

Venkata S Govindarajan, Katherine Atwell, Barea Sinno, Malihe Alikhani, David Beaver, and Junyi Jessy Li. How people talk about each other: Modeling generalized intergroup bias and emotion. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 2488–2498, Dubrovnik, Croatia, May 2023a. Association for Computational Linguistics.

Venkata S Govindarajan, David Beaver, Kyle Mahowald, and Junyi Jessy Li. Counterfactual probing for the influence of affect and specificity on intergroup bias. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 12853–12862, Toronto, Canada, July 2023b. Association for Computational Linguistics.

Anthony G. Greenwald and Linda Hamilton Krieger. Implicit bias: Scientific foundations. 94(4):945–967. ISSN 0008-1221. Publisher: California Law Review, Inc.

Dirk Groeneveld, Iz Beltagy, Pete Walsh, Akshita Bhagia, Rodney Kinney, and et al. Olmo: Accelerating the science of language models. *Preprint*, 2024.

Lauren Hall-Lew, Emma Moore, and Robert J. Podesva. *Social Meaning and Linguistic Variation: Theoretical Foundations*, page 1–24. Cambridge University Press, 2021.

David Heineman, Yao Dou, and Wei Xu. Thresh: A unified, customizable and deployable platform for fine-grained text evaluation. In Yansong Feng and Els Lefever, editors, *Proceedings of the 2023 Conference on Empirical Methods in*

Natural Language Processing: System Demonstrations, pages 336–345, Singapore, December 2023. Association for Computational Linguistics.

W. Hoppel, Denise Sekaquaptewa, and P. Vargas. The Linguistic Intergroup Bias As an Implicit Indicator of Prejudice. *Journal of Experimental Social Psychology*, 33:490–509, 1997.

Maksim Horowitz, Ron Yurko, and S Ventura. nflscraper: Compiling the nfl play-by-play api for easy use in r, 2017. URL <https://github.com/maksimhorowitz/nflscraper>.

Shanto Iyengar, Yphtach Lelkes, Matthew Levendusky, Neil Malhotra, and Sean J. Westwood. The origins and consequences of affective polarization in the united states. 22(1):129–146.

Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, and Devendra Singh Chaplot et al. Mistral 7b, 2023.

Masahiro Kaneko and Danushka Bollegala. Gender-preserving Debiasing for Pre-trained Word Embeddings. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 1641–1650, Florence, Italy, July 2019. Association for Computational Linguistics.

Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.

Jayant Krishnamurthy and Thomas Kollar. Jointly learning to parse and perceive: Connecting natural language to the physical world. *Transactions of the Association for Computational Linguistics*, 1:193–206, 2013.

Junyi Jessy Li. *From Discourse Structure To Text Specificity: Studies Of Coherence Preferences*. phdthesis, University of Pennsylvania, 2017.

Percy Liang, Michael Jordan, and Dan Klein. Learning semantic correspondences with less supervision. In Keh-Yih Su, Jian Su, Janyce Wiebe, and Haizhou Li, editors, *Proceedings of the Joint Conference of the 47th Annual Meeting of the ACL and the 4th International Joint Conference on Natural Language Processing of the AFNLP*, pages 91–99, Suntec, Singapore, August 2009. Association for Computational Linguistics.

Annie Louis and Ani Nenkova. Text specificity and impact on quality of news summaries. In *Proceedings of the Workshop on Monolingual Text-To-Text Generation*, pages 34–42, Portland, Oregon, June 2011. Association for Computational Linguistics.

Anne Maass. Linguistic Intergroup Bias: Stereotype Perpetuation Through Language. In Mark P. Zanna, editor, *Advances in Experimental Social Psychology*, volume 31, pages 79–121. Academic Press, 1999.

Anne Maass, Daniel Anthony Salvi, Luciano Arcuri, and Gün R. Semin. Language use in intergroup contexts: the linguistic intergroup bias. *Journal of personality and social psychology*, 57 6:981–993, 1989.

Jack Merullo, Luke Yeh, Abram Handler, Alvin Grissom II, Brendan O’Connor, and Mohit Iyyer. Investigating sports commentator bias within a large corpus of American football broadcasts. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 6355–6361, Hong Kong, China, November 2019. Association for Computational Linguistics.

Saif Mohammad. #Emotional Tweets. In **SEM 2012: The First Joint Conference on Lexical and Computational Semantics – Volume 1: Proceedings of the*

main conference and the shared task, and Volume 2: Proceedings of the Sixth International Workshop on Semantic Evaluation (SemEval 2012), pages 246–255, Montréal, Canada, 7-8 June 2012. Association for Computational Linguistics.

Saif M. Mohammad and Svetlana Kiritchenko. Using Hashtags to Capture Fine Emotion Categories from Tweets. *Computational Intelligence*, 31:301 – 326, 2015.

Saif M. Mohammad and Peter D. Turney. Crowdsourcing a Word-Emotion Association Lexicon. *Computational Intelligence*, 29, 2013.

Courtney Napoles, Keisuke Sakaguchi, Matt Post, and Joel Tetreault. Ground truth for grammatical error correction metrics. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 588–593, Beijing, China, July 2015. Association for Computational Linguistics.

Dat Quoc Nguyen, Thanh Vu, and Anh Tuan Nguyen. BERTweet: A pre-trained language model for English Tweets. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 9–14. Association for Computational Linguistics, 2020.

Konstantinos Pelechrinis and Evangelos Papalexakis. The anatomy of american football: evidence from 7 years of nfl game data. *PLoS one*, 11(12):e0168716, 2016.

Robert Plutchik. The Nature of Emotions. *American Scientist*, 89(4):344–350, 2001.

Reid Pryzant, Richard Diehl Martinez, Nathan Dass, Sadao Kurohashi, Dan Jurafsky, and Diyi Yang. Automatically Neutralizing Subjective Bias in Text.

Proceedings of the AAAI Conference on Artificial Intelligence, 34(01):480–489, Apr. 2020.

Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*, 2021.

Shauli Ravfogel, Yanai Elazar, Hila Gonen, Michael Twiton, and Yoav Goldberg. Null it out: Guarding protected attributes by iterative nullspace projection. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7237–7256, Online, July 2020. Association for Computational Linguistics.

Shauli Ravfogel, Grusha Prasad, Tal Linzen, and Yoav Goldberg. Counterfactual interventions reveal the causal effect of relative clause representations on agreement prediction. In *Proceedings of the 25th Conference on Computational Natural Language Learning*, pages 194–209, Online, November 2021. Association for Computational Linguistics.

Tim Sainburg, Leland McInnes, and Timothy Q Gentner. Parametric UMAP Embeddings for Representation and Semisupervised Learning. *Neural Computation*, 33(11):2881–2907, 2021.

Maarten Sap, Saadia Gabriel, Lianhui Qin, Dan Jurafsky, Noah A. Smith, and Yejin Choi. Social bias frames: Reasoning about social and power implications of language. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5477–5490, Online, July 2020. Association for Computational Linguistics.

Sherry B Schnake and Janet B Ruscher. Modern racism as a predictor of the linguistic intergroup bias. *Journal of Language and Social Psychology*, 17(4): 484-491, 1998.

G. R. Semin and K. Fiedler. The cognitive functions of linguistic categories in describing persons: Social cognition and language. 54:558-568, 1988. ISSN 0022-3514. Publisher: American Psychological Association.

Emily Sheng, Kai-Wei Chang, Premkumar Natarajan, and Nanyun Peng. The Woman Worked as a Babysitter: On Biases in Language Generation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3407-3412, Hong Kong, China, November 2019. Association for Computational Linguistics.

Emily Sheng, Kai-Wei Chang, Prem Natarajan, and Nanyun Peng. Towards Controllable Biases in Language Generation. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3239-3254, Online, November 2020. Association for Computational Linguistics.

Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, and et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *Transactions on Machine Learning Research*, 2023. ISSN 2835-8856.

Gemini Team. Gemini: A family of highly capable multimodal models, 2023.

Teun A Van Dijk. *Society and Discourse: How Social Contexts Influence Text and Talk*. Cambridge University Press, 2009.

Achille Varzi. Mereology. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Spring 2019 edition, 2019.

Somin Wadhwa, Silvio Amir, and Byron C. Wallace. Revisiting relation extraction in the era of large language models. *Proceedings of the conference. Association for Computational Linguistics. Meeting*, 2023:15566–15589, 2023.

Sida Wang and Christopher Manning. Baselines and Bigrams: Simple, Good Sentiment and Topic Classification. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 90–94, Jeju Island, Korea, July 2012. Association for Computational Linguistics.

Wenbo Wang, Lu Chen, Krishnaprasad Thirunarayan, and Amit P. Sheth. Harnessing Twitter ”Big Data” for Automatic Emotion Identification. In *2012 International Conference on Privacy, Security, Risk and Trust and 2012 International Conference on Social Computing*, pages 587–592, 2012.

Albert Webson, Zhizhong Chen, Carsten Eickhoff, and Ellie Pavlick. Are “Undocumented Workers” the Same as “Illegal Aliens”? Disentangling Denotation and Connotation in Vector Spaces. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4090–4105, Online, November 2020. Association for Computational Linguistics.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, brian ichter, Fei Xia, Ed Chi, Quoc V Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems*, volume 35, pages 24824–24837. Curran Associates, Inc., 2022.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien

Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander M. Rush. Transformers: State-of-the-Art Natural Language Processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45, Online, October 2020. Association for Computational Linguistics.

J.P. Woodard and J.T. Nelson. An information theoretic measure of speech recognition performance. 1982.

Ronald Yurko, Samuel L. Ventura, and Maksim Horowitz. nflwar: a reproducible method for offensive player evaluation in football. *Journal of Quantitative Analysis in Sports*, 15:163 – 183, 2018.

Samira Zad, Joshuan Jimenez, and Mark Finlayson. Hell Hath No Fury? Correcting Bias in the NRC Emotion Lexicon. In *Proceedings of the 5th Workshop on Online Abuse and Harms (WOAH 2021)*, pages 102–113, Online, August 2021. Association for Computational Linguistics.